

Overlapping Generations in Lean: The Life-Cycle Household

LeanEconomics Org [mostly Claude]*

September 2026

Draft, in progress

Abstract

This is the record of a programme to formalise overlapping-generations economies in Lean 4. In Diamond's two-period economy everything is closed form and the steady state is unique for every relative risk aversion, with no condition on the savings function, even though the saving rate itself falls with the return whenever risk aversion exceeds one: the wage channel of a Cobb-Douglas technology dominates. That is the only place in the paper where uniqueness is settled outright, and it shows that a proof need not go through the household-side comparative static at all. A household that lives J periods against uninsurable earnings risk and a borrowing constraint solves a finite backward induction, so everything proved for the infinite-horizon Bellman operator with an arbitrary continuation applies stage by stage with no fixed point. On top of that we prove the finite-horizon sandwich, $\kappa_k m \leq c_k \leq \kappa_k(m + H_k(z))$ with $1/\kappa_k = \sum_{i \leq k} (b/R)^i$, and then that κ_k is exactly the reciprocal of that geometric sum, so the gap to its infinite-horizon limit closes at a known rate. Adding retirement and a pension, we settle a claim often made informally: a means test does make the policy non-monotone, but not for the reason usually given. Saving rises with cash on hand whatever the continuation, by Topkis, so a value function made non-concave by the test cannot by itself reverse it; what the test reverses is consumption, and cash on hand itself once the taper is steeper than the gross return. The main result is existence of a stationary equilibrium with both boundary conditions discharged: capital supply is a finite composition, with no stationary distribution, no Feller property and no Doeblin condition, and it is bracketed above and below by affine bounds on a cohort's accumulated assets whose coefficients are explicit recursions in the discount factor, the interest rate, the earnings process and the horizon. Nothing in the hypotheses refers to the household's decisions, and at Aiyagari's calibration they hold on an interval of five percentage points containing the true equilibrium. We also report what the life cycle does for Light's Theorem 1, that saving rises with the interest rate: it holds at every age below a critical relative risk aversion of about 3.4, so the life cycle delivers Aiyagari's $\mu = 3$ and not his $\mu = 5$; at logarithmic utility we give

*The Lean development described here is at <https://github.com/LeanEconomics/LeanEconomics>. It was written mostly by Claude (Anthropic), an AI coding model, working under the direction of Robert Kirkby (Victoria University of Wellington); every result reported was checked by the Lean kernel. This is the companion to the write-up on Aiyagari (1994) in the same repository. Comments welcome.

a sufficient condition that is true where the Euler condition is false. Where the bounds stop is the same obstruction at both ends, that a power mean with negative exponent is superadditive rather than 1-Lipschitz.

JEL: C63, D52, E21. *Keywords:* overlapping generations, life cycle, Bewley models, formal verification, Lean.

1 Introduction

The companion write-up in this repository proved that the stationary equilibrium of Aiyagari (1994) is unique at his logarithmic calibration, and located exactly what stands in the way for relative risk aversion $\mu \in \{3, 5\}$. Every obstacle it met traces to the fixed point of the infinite-horizon problem: a condition on the response of consumption to the interest rate fails in a tail of wealth that the stationary distribution reaches; an induction that would prove the response bound has a step multiplier $1/p > 1$, so it cannot close as a supremum argument; and the only floor on capital supply comes from an ergodic identity that loses a factor $\|u'\|^2 / \text{Var}$. Overlapping generations remove all three at once. A household that lives J periods solves a finite backward induction from an explicit terminal condition, so a multiplier $1/p$ per step compounds to p^{-J} , which at Aiyagari's numbers is about 1.03. Wealth is bounded by J periods of accumulation, so any condition needs checking on a known bounded region. And the distribution of the youngest cohort is common to every interest rate, so an ordering of policies pushes forward to an ordering of the whole age profile of distributions by induction on age, with no Doeblin atom and no uniqueness of a stationary distribution.

The target of the programme is uniqueness of the stationary equilibrium of a stochastic life-cycle economy in the manner of Huggett (1996) at $\mu \in \{3, 5\}$. That target is not reached here, and the paper is as much a record of where it is blocked as of what is proved.

What is proved is this. Section 2 does Diamond's two-period economy, where the steady state is unique for every risk aversion and the household-side comparative static that blocks the general case runs the wrong way without doing any harm. Section 3 sets out the life-cycle household and the finite-horizon sandwich on consumption. Section 5 adds retirement, a pension and a means test, and settles which policy a means test makes non-monotone and in which state space. Section 6 is the main result: a stationary equilibrium exists, with both boundary conditions discharged from primitives rather than assumed. Capital supply in a life cycle is a finite composition of finite sums, so existence needs no stationary distribution and no compactness argument beyond Berge; and it can be bracketed by two affine bounds on a cohort's accumulated assets, one built from a floor on saving and one from a ceiling, sharing a slope, with every coefficient an explicit recursion in the primitives. At Aiyagari's calibration the hypotheses hold on an interval of about five percentage points containing the true equilibrium rate of 5.09%.

What is not proved is Light's Theorem 1, that saving rises with the interest rate, which is the one household-side input uniqueness still needs. Section 7 reports how far it goes: it is

true at every age up to a critical relative risk aversion of about 3.4, so the life cycle covers Aiyagari's $\mu = 3$ and not his $\mu = 5$, and the Euler reduction of it is false at logarithmic utility exactly where an equilibrium sits. A different sufficient condition, on the value gap between two rates, is true there instead, and Section 8 proves it suffices. Section 11 records the rest of the programme, the classical results of Diamond (1965) included, and what each stage needs.

As in the companion, every result stated with a formal name has been checked by the Lean kernel against Mathlib; the development is at <https://github.com/LeanEconomics/LeanEconomics>, directory `LeanEconomics/OLG`, and the numerical work is in `WriteUps/WriteUpOLG/numerics`.

2 Diamond's two-period economy

The base case of the programme, and the one place in it where uniqueness is proved outright rather than reduced to an assumption. Agents live two periods, work when young and consume out of savings when old; a representative firm rents capital from the old.

2.1 The household, in closed form

A young household with wage w facing gross return R solves $\max u(w - s) + \beta u(Rs)$. With CRRA utility the Euler equation $u'(w - s) = \beta R u'(Rs)$ has the explicit solution $s = \sigma(R)w$ with

$$\sigma(R) = \frac{1}{1 + \beta^{-1/\gamma} R^{(\gamma-1)/\gamma}}$$

(`saveShare`). That σ satisfies the Euler equation is algebra (`saveShare_euler`); that it is the optimum needs one more ingredient.

Proposition 2.1 (CRRA lies below its tangent). *For $c, c' > 0$ and $\gamma > 0$, $u(c') \leq u(c) + u'(c)(c' - c)$.*

Formal statement: `crra_le_tangent`.

This is Bernoulli's inequality in both directions. Below unit risk aversion the exponent $1 - \gamma$ lies in $(0, 1)$ and the concave form applies directly; above it the exponent is negative, and inverting the consumption ratio and multiplying through turns the claim into the form with exponent $\gamma \geq 1$. Optimality then follows by adding the tangent inequality at the two consumption levels and cancelling with the Euler equation, with no differentiation of a value function anywhere (`saveShare_isOptimal`).

2.2 The comparative static is discouraging

Proposition 2.2 (The textbook comparative static). *The saving rate rises with the gross return if $\gamma < 1$ and falls if $\gamma > 1$.*

Formal statement: `saveShare_strictMonoOn`.

, `saveShare_strictAntiOn`

So the household-side monotonicity that Section 7 cannot establish in general is here a one-line calculation, and it runs the *wrong* way at the risk aversions economists use: above unit intertemporal elasticity the income effect beats the substitution effect and a household facing a better return saves a smaller share of a given wage.

2.3 The steady state is unique anyway

What rescues it is general equilibrium. With Cobb-Douglas production the wage rises and the return falls as capital deepens, and the wage channel dominates. Writing $k' = \sigma(R(k)) w(k)$ for capital per worker next period (`diamondMap`), the whole argument is that k'/k falls.

Proposition 2.3 (The engine). *If $\theta \leq 1$, $B > 0$, $a \geq 0$ and $b > 0$, then $x \mapsto x/(1+B(a+bx)^\theta)$ is strictly increasing on $(0, \infty)$.*

Formal statement: `div_one_add_rpow_strictMono`.

The proof is the chain $x_2/x_1 \leq R_2/R_1 \leq (R_2/R_1)^\theta$, where the first step is $ax_2 \leq ax_1$ and the second is that a ratio at most one rises when raised to a power at most one. No calculus. Substituting $x = k^{\alpha-1}$, which falls with k , and $\theta = (\gamma - 1)/\gamma$, which is below one for every positive γ , gives the result.

Theorem 2.4 (Existence and uniqueness). *For every $\gamma > 0$, every $\delta \leq 1$ and Cobb-Douglas production, the map $k \mapsto k'/k$ is strictly decreasing, so there is at most one positive steady state; and if the map crosses the diagonal between two positive levels there is exactly one.*

Formal statement: `exists_unique_steadyState`.

, `steadyState_unique`, `exists_steadyState`

No condition on the savings function is needed, of the kind textbook treatments impose. The numerical check is in `numerics/dcheck.m`: the closed form matches a grid search to within the grid spacing at every risk aversion and return tried; there is exactly one positive crossing of the diagonal for risk aversions from 0.5 to 10 at both $\delta = 1$ and $\delta = 1/2$; and the saving rate rises with the return only below unit risk aversion, as Proposition 2.2 says.

What this does and does not settle is worth stating. It settles uniqueness in the two-period economy completely, for every risk aversion, which the stochastic life cycle of Section 6 does not. It settles it by an argument that has no counterpart there, because with more than two periods the household's saving is not a closed form and k'/k is not a product of two monotone factors. What it does supply is the shape of the answer: the household-side comparative static can fail while the general-equilibrium one holds, so a proof of uniqueness need not go through Light's Theorem 1 at all.

3 The life-cycle household

3.1 Stages

The household of the companion write-up is kept: earnings $y(z)$ follow a finite Markov chain π , the gross return is $R = 1 + r$, cash on hand is $m = Ra + y(z)$, saving $a' \in [0, \bar{a}]$ is

bounded below by the borrowing limit and above by an artefactual cap, utility is CRRA with $u'(c) = c^{-\mu}$ and $\beta \in (0, 1)$. What changes is the horizon: the household lives J periods and leaves nothing behind. Count periods from the end. With k periods still to come after today, the continuation value is the k -th Bellman iterate from zero,

$$V_k = T^k 0, \quad (Tv)(a, z) = \max_{a'} \{u(m - a') + \beta \sum_{z'} \pi(z, z') v(a', z')\},$$

and the household with k periods to come saves $g_k = \text{policyOf } V_k$ and consumes $c_k = m - g_k$ (`stageValue`, `stagePolicy`, `stageConsumption`). Age j of a J -period life is stage $k = J - 1 - j$. In the last period the household saves nothing and consumes its cash on hand (`stagePolicy_zero`, `stageConsumption_zero`).

This is the finite-horizon problem with a flat age profile of earnings: the same income process at every age. An age profile, retirement and survival risk need income to be a parameter of the Bellman step rather than a field of the structure, and are deferred to the next stage of the programme.

3.2 What transfers without a fixed point

The companion development proved its household results for the Bellman operator applied to an arbitrary continuation v , and only then passed to the fixed point. Here there is no fixed point to pass to, and each result is read at $v = V_k$:

- consumption is positive at every stage when marginal utility is unbounded at zero (`stageConsumption_pos`);
- the stage value has concave slices, so the optimal action is unique and both saving and consumption rise with wealth (`stagePolicy_mono`, `stageConsumption_mono`);
- Carroll–Kimball: consumption is concave in wealth at every stage (`concaveOn_stageConsumption`);
- the Euler inequality under the cap, $\beta R \mathbb{E}[u'(c_k(a', z')) \mid z] \leq u'(c_{k+1}(a, z))$, at every state (`stage_euler_le`), and the Euler inequality above the floor, the reverse, at every state where the household saves (`stage_euler_ge`).

3.3 The finite-horizon sandwich

Let $\mathfrak{p} = (\beta R)^{1/\mu}$ and define the finite-horizon minimal propensities to consume by

$$\kappa_0 = 1, \quad \kappa_{k+1} = \frac{\kappa_k R}{\mathfrak{p} + \kappa_k R}, \quad \text{so that} \quad \frac{1}{\kappa_k} = \sum_{i=0}^k \left(\frac{\mathfrak{p}}{R}\right)^i$$

(`stageMPC`): the share of cash on hand consumed by a household with k periods to come and no future income. Let $H_0 = 0$ and $RH_{k+1}(z) = \sum_{z'} \pi(z, z')(y(z') + H_k(z'))$, the expected present value of the next k incomes (`stageHumanWealth`).

Theorem 3.1 (The finite-horizon sandwich). *Assume CRRA utility with marginal utility unbounded at zero, and that the cap is slack at every stage. Then for every k , every z and every $a \in [0, \bar{a}]$,*

$$\kappa_k m \leq c_k(a, z) \leq \kappa_k(m + H_k(z)).$$

Formal statement: *IncomeFluctuation.stageMPC_mul_le_stageConsumption, IncomeFluctuation.stageConsumption_le_stageHumanWealth.*

Each bound is one Euler step, applied k times. The lower bound: if consumption against v is at least κm , the Euler inequality above the floor at a state where the household saves gives $c^{-\mu} \leq \beta R \mathbb{E}[(\kappa R a')^{-\mu}]$ with $a' = m - c$, hence $\kappa R(m - c) \leq \mathfrak{p}c$ and $c \geq \kappa R(\mathfrak{p} + \kappa R)^{-1}m$; at a corner $c = m$ (*ccrra_stageMPC_step*). The upper bound: if consumption against v is at most $\kappa(m + H(z))$, the Euler inequality under the cap and Jensen for the negative power give $c^{-\mu} \geq \beta R (\mathbb{E}[c'])^{-\mu}$ with $\mathbb{E}[c'] \leq \kappa R(a' + H'(z))$, hence $c \leq \kappa R(\mathfrak{p} + \kappa R)^{-1}(m + H'(z))$ (*ccrra_stageUpper_step*). The recursion for κ is the same on both sides, which is why the sandwich is exact in the homogeneous case.

The propensities decrease in the horizon towards the infinite-horizon $\kappa = 1 - \mathfrak{p}/R$ of the companion write-up: $\kappa \leq \kappa_k$ for every k (*minMPC_le_stageMPC*), so a household with less life left consumes a larger share of its cash, and the infinite-horizon bounds are the limiting case. A first consequence for the programme is an explicit bound on saving, $g_k(a, z) \leq (1 - \kappa_k)(Ra + y_{\max})$ (*stagePolicy_le*), which iterated from zero wealth at the start of life bounds the wealth of every cohort.

4 The deterministic life cycle

Between Diamond's two periods and the stochastic life cycle sits the case with many periods and no earnings risk, the Auerbach–Kotlikoff class. It is worth doing on its own, because there the sandwich of Theorem 3.1 closes and Light's Theorem 1 becomes algebra.

4.1 The sandwich closes

Make income the same in every state, and suppose the household is interior at every age: it saves something and stops short of the artefactual cap.

Theorem 4.1 (Consumption in closed form). *Under those two conditions, $c_k(a, z) = \kappa_k(m + H_k)$ exactly, so saving is exactly $(1 - \kappa_k)(y + Ra) - \kappa_k H_k$.*

Formal statement: *IncomeFluctuation.stageConsumption_eq_of_interior.*

The proof squeezes the Euler equation between its two directions. Interiority gives both the inequality above the floor and the one under the cap, hence equality; constant income makes tomorrow's consumption the same in every state, so the expectation collapses to a point and the power can be inverted; the budget identity closes the induction. What the stochastic model can only bracket is here an identity.

4.2 Uniqueness at logarithmic utility

Theorem 4.2 (Theorem 1, deterministic and logarithmic). *At $\gamma = 1$, saving rises with the interest rate at every age and every asset level, with no condition. Aggregate capital rises with it, and the stationary equilibrium rate is unique.*

Formal statement: `IncomeFluctuation.stagePolicy_mono_det_log.`

`, IncomeFluctuation.olgCapital_mono_det_log`

Two facts do it. At $\gamma = 1$ the patience factor is βR , so the R cancels out of the recursion for κ and the propensity does not depend on the rate at all (`stageMPC_one_eq_of_rate`); and human wealth falls as the rate rises (`stageHumanWealth_antitone_rate`). The gap is then exactly

$$g(R_2) - g(R_1) = (1 - \kappa_k)(R_2 - R_1)a + \kappa_k(H_k(R_1) - H_k(R_2)),$$

and both terms are nonnegative. Compare Section 8, where at logarithmic utility with risk the Euler condition is *false* at the rates that matter and a different sufficient condition has to be found. Risk is what makes it hard, not the horizon.

4.3 Above logarithmic utility, and a correction

For $\gamma \neq 1$ the propensity moves with the rate and the comparative static becomes a duration comparison. Consumption at zero assets is $y \sum_{i \leq K} R^{-i} / \sum_{i \leq K} q^i$ with $q = \beta/R$, and it falls with the rate exactly when

$$T_y \geq \left(1 - \frac{1}{\gamma}\right) T_c,$$

where T_y and T_c are the durations of the income and consumption profiles. Positive assets only help, since they add $(1 - \kappa_k)(R_2 - R_1)a \geq 0$.

It is worth being careful about what that threshold is. The certainty-equivalent analysis of Section 7 reports a critical risk aversion of 4.33, and it is tempting to read that as the deterministic threshold. It is not. That number is the certainty equivalent of the *risky* problem read at the state that matters, where expected income falls back towards the mean and so is front-loaded relative to consumption. With a flat income profile the two durations are nearly equal, and the criterion has enormous room (`numerics/detcheck.m`):

r	0%	2%	4%	6%
Income duration T_y	29.50	23.70	18.70	14.79
Consumption duration T_c	28.89	23.40	18.68	14.95
T_y/T_c	1.021	1.013	1.001	0.989
Criterion binds at γ of	never	never	never	94

Below the rate of time preference the income profile is in fact the *longer* of the two, so $T_y > T_c$ and the criterion holds for every γ whatever; above it consumption is the longer, and the criterion first fails at a risk aversion near 94. Either way there is no threshold an economist would meet: the derivative of saving with respect to the rate at $r = 4\%$ is positive and comfortably so from $\gamma = 1$ to $\gamma = 5$, at zero assets and at five units of them. The

threshold of Section 7 is a feature of the earnings *profile* the risky problem induces, not of risk aversion by itself. Only the logarithmic case is formalised, because above it the duration comparison is an inequality between two geometric sums with different ratios, which is true here but not unconditionally.

5 Retirement, a pension, and the means test

5.1 Retirement with a lump-sum pension

The household works until age J_r and is retired after it, receiving a pension in place of its labour earnings. With a *lump-sum* pension p the retired household faces the same problem as the working one with income replaced by the constant p , so the retired phase is itself an income-fluctuation problem (`withPension`) and everything in Section 3 applies to it unchanged. The life-cycle value is assembled by iterating the working Bellman operator on the retired value: $\text{lifeValue}(n_r, n_w)$ is the value of a household with n_w working and n_r retired periods still to come (`lifeValue`), so a household of age $j < J_r$ in a life of J periods sits at $n_w = J_r - j$ and $n_r = J - J_r$. Saving and consumption rise with wealth at every age (`lifePolicy_mono`, `lifeConsumption_mono`).

Two facts about the pension itself are worth recording. A retiree's human wealth is the same in every earnings state, because the pension does not depend on the earnings shock (`stageHumanWealth_withPension_const`), and it never exceeds the perpetuity value of the pension, $H_k \leq p/r$ (`stageHumanWealth_withPension_le`). With Theorem 3.1 the retiree's consumption therefore satisfies

$$\kappa_k m \leq c_k(a) \leq \kappa_k \left(m + \frac{p}{r} \right)$$

at every retired age. And cash on hand is strictly increasing in assets (`withPension_resources_strictMono`). Both of these fail once the pension is means-tested.

5.2 Means testing

A pension is means-tested when the amount paid falls with the recipient's assets. We write the schedule as a function $p(\cdot)$ of assets and cash on hand as $m(a) = Ra + p(a)$ (`cash`), and consider three schedules (`lumpSum`, `tapered`, `cliff`): a lump sum; a taper, withdrawn at rate τ per unit of assets above a threshold and never negative; and a cliff, withdrawn entirely above the threshold.

It is often argued that a means test makes the household's policy function non-monotone, and the reason usually given is that the test makes the value function non-concave. The argument as usually given is wrong, and the conclusion is nevertheless true, for a different reason. Three results separate the two.

Theorem 5.1 (Monotonicity in cash on hand is never lost). *Let period utility be strictly concave on the positive half-line and let W be any continuation value. If $m_1 < m_2$ and x_i is optimal at m_i , then $x_1 \leq x_2$.*

Formal statement: *OLG.isOptimalSaving_mono.*

The proof is Topkis. Adding the two optimality inequalities gives $u(m_1 - x_1) + u(m_2 - x_2) \geq u(m_1 - x_2) + u(m_2 - x_1)$; if $x_2 < x_1$ the two right-hand consumptions are strict convex combinations of the two left-hand ones with the same sum, so strict concavity makes the reverse inequality strict, a contradiction. The continuation enters only through terms that cancel. So a non-concave value function — which a means test certainly produces — cannot by itself make saving fall with cash on hand, whatever the shape of W .

What a non-concave continuation does produce is a jump, and the jump falls on consumption.

Theorem 5.2 (Consumption falls when wealth rises). *There are a continuation value arising from a means-tested pension and two levels of cash on hand, the second larger, at which optimal consumption is strictly smaller. With log utility, a unit gross return, and a pension of one withdrawn entirely above assets of one: at $m = 6$ the household saves exactly 1 and consumes 5; at $m = 8$ it saves 4 and consumes 4.*

Formal statement: *OLG.exists_consumption_decrease.*

At $m = 6$ the household keeps the pension, saving the most it can without losing it; at $m = 8$ the pension is no longer worth keeping, so the household gives it up and saves the proceeds. Saving rises from 1 to 4, as Theorem 5.1 says it must, and since it rises by more than wealth does, consumption falls. Both optima are exact: the objective is log of a quadratic on each side of the threshold, and the comparison is checked by the kernel.

The third result is the genuine non-monotonicity of the saving policy, and it comes from the budget rather than from the value function. A retiree's pension is assessed each period on the assets then held, so the schedule enters cash on hand directly.

Theorem 5.3 (A gentle taper is harmless; a cliff is not). *If $0 \leq \tau < R$ then $a \mapsto Ra + p(a)$ is strictly increasing under the taper (*cash_strictMono_tapered*), and so, by Theorem 5.1, saving is monotone in assets (*isOptimalSaving_mono_assets*). Under a cliff it is not: if $a \leq \bar{a} < b$ and $R(b - a) < p$ then cash on hand is strictly smaller at b (*cash_lt_cash_of_cliff*), and there are two asset levels, the second larger, at which the household saves strictly less: with a concave continuation, at assets 1 it saves $1/2$ and at assets $3/2$ only $1/4$.*

Formal statement: *OLG.exists_saving_decrease_in_assets.*

The continuation in the second half of Theorem 5.3 is $\log(x+1)$, which is concave: nothing is hidden in it. The non-monotonicity is entirely the budget's, and the threshold is sharp, τ against R . A taper gentler than the gross return leaves the household's saving monotone in assets no matter how badly the test distorts the value function; a taper steeper than the gross return, or a cliff, makes a richer household poorer in the only sense that matters for today's decision.

5.3 Which policy, and in which variable

The two non-monotonicity results are about different objects, and it is worth separating them in the household's own state space, assets a and the Markov income shock z .

Next period's assets are monotone. For each z , the optimal a' is nondecreasing in a whatever the continuation value, with no concavity assumed of it:

Theorem 5.4 (a' rises with a , at each z). *Let $a < a'$ both lie in the asset region. Then for every continuation v and every z , $g_v(a, z) \leq g_v(a', z)$.*

Formal statement: *IncomeFluctuation.policyOf_mono_of_lt.*

This is Theorem 5.1 read in the state space: cash on hand rises with assets, so optimal saving does. The repository already had this statement (`policyOf_mono`) but with concave slices of the continuation as a hypothesis, which is exactly what a means test destroys; the Topkis argument removes the hypothesis, at the cost of requiring the two asset levels to be distinct. The continuation cancels between the two optimality inequalities and only strict concavity of period utility survives.

Consumption is not monotone, and the contrast is sharp rather than an artefact of proof technique. Consumption is $m(a) - g_v(a, z)$, and both terms rise with a ; when the continuation is non-concave the second can rise faster, which is exactly Theorem 5.2. So monotone consumption genuinely requires concavity of the continuation, while monotone saving does not.

Two things follow for the means test. Where the pension enters only next period's problem, so that today's cash on hand is unaffected, Theorem 5.4 applies verbatim: next period's assets still rise with this period's, and only consumption misbehaves. Where the test is assessed on the assets currently held, the hypothesis of Theorem 5.4 is the one that fails: cash on hand is no longer increasing in a , and with it the policy. That is the whole of Theorem 5.3, and the dividing line is τ against R .

5.4 At realistic parameters

The three theorems predict exactly which of the two channels operates for which schedule, and a calibrated retiree confirms it. We solve a deterministic retiree of 25 years by backward induction with global search over a grid of 6001 asset levels — the value function is not concave, so a first-order condition cannot be used — at $r = 4\%$, $\beta = 0.96$, relative risk aversion 3, a full pension of 0.40 of mean earnings and a free area of 3 mean-earnings units (`numerics/meanstest.m`). The taper of 0.078 per unit of assets is Australia's asset test, three dollars a fortnight per thousand dollars.

Pension schedule	Cash on hand	Saving	Consumption
Lump sum	increasing	increasing	increasing
Taper 0.078 (Australia)	increasing	increasing	falls by 0.025–0.046
Taper 0.50	increasing	increasing	falls by 0.13–0.24
Taper 1.20 (above R)	falls	falls by 0.004	falls by 0.55–0.70
Cliff	falls by 0.40	falls by 0.32–0.37	falls by 0.55–0.70

Every entry is what the theorems require. Saving falls with assets in exactly the two rows where cash on hand does, which are exactly the two rows whose taper exceeds the

gross return, as Theorem 5.3 says; in the other three it rises, as Theorem 5.1 says it must. Consumption falls in every row with a test of any kind, as Theorem 5.2 says it can, and the drop sits exactly at the free area and grows with age, from 0.025 of mean earnings at the start of retirement to 0.046 twenty years in under the Australian taper. The mechanism is the one the witness isolates: just above the free area the effective return to saving is only $R - \tau$ until the pension is exhausted, so the household jumps its saving toward the far side of the taper range and cuts consumption to pay for it.

The effect is not confined to retirees. A worker in the last year before retirement, facing the Australian taper for the whole of retirement, has a monotone saving policy and a consumption policy that falls by 0.025 of mean earnings at the same threshold: the non-concavity is inherited through the continuation value, and it is again consumption that carries it.

6 Stationary equilibrium

Every cohort is born with no assets and an earnings state drawn from a fixed law ν , lives J periods, and is replaced; cohorts have equal weight. The distribution of the youngest cohort is therefore the same at every interest rate, and the distribution of every older cohort is that law pushed forward by the policies of the ages in between. This is the structural advantage of the life cycle over the infinite horizon, and it removes most of the machinery the stationary case needed.

Two of the consumption bounds this section uses are developed in Section 7, where they arise in the hunt for the critical risk aversion: the constraint-incidence bound of Theorem 7.8 and the asset-dependent saving floor of Section 7.13. Nothing here depends on the verdict of that section.

6.1 The aggregate is a finite composition

Because assets at birth are degenerate at zero, no measure theory is required. Let

$$(\mathcal{S}_k h)(a, z) = \sum_{z'} \pi(z, z') h(g_k(a, z), z')$$

be the expectation of h next period for a household with k periods still to come (`stageStep`). The assets a household holds over the rest of its life, summed across ages and seen from its current state, satisfy

$$A_0(a, z) = a, \quad A_{k+1}(a, z) = a + (\mathcal{S}_{k+1} A_k)(a, z)$$

(`cohortAssets`), and aggregate capital per head is $K = \frac{1}{J} \sum_z \nu(z) A_{J-1}(0, z)$ (`olgCapital`). There is no pushforward measure, no Feller property, no stationary distribution, no Doeblin condition and no uniqueness of an invariant law: the equilibrium object is a finite sum of finite sums, and the recursion above is its definition.

6.2 The chain

Theorem 6.1 (Light’s Theorem 2 for the life cycle). *Suppose two economies share the earnings chain and their policies are ordered at every age, $g_k^{(1)} \leq g_k^{(2)}$ on the asset region. Then $A_k^{(1)} \leq A_k^{(2)}$ for every k , and aggregate capital is ordered, $K^{(1)} \leq K^{(2)}$.*

Formal statement: *IncomeFluctuation.cohortAssets_le, IncomeFluctuation.olgCapital_le_of_stagePolicy_le.*

The proof is an induction on age with two ingredients, and both are already available. The class of functions increasing in assets is preserved by $\mathcal{S}_k$ (`monoRegion_stageStep`), because next period’s assets rise with this period’s at each earnings state and for any continuation, which is Theorem 5.4. And ordered policies keep ordered functions ordered (`stageStep_le_stageStep`): the higher-rate economy evaluates a larger function at a larger asset level. Note what is *not* needed. The infinite-horizon version of this step required monotone transitions, a coupling argument and first-order stochastic dominance of the stationary distributions; here the whole statement is an inequality between two functions on the same finite state space, because the distribution the two economies start from is literally the same.

6.3 Uniqueness

Theorem 6.2 (Single crossing). *A capital supply that never falls with the rate meets a strictly falling demand at most once.*

Formal statement: *IncomeFluctuation.olgEquilibriumRate_unique.*

With Theorem 6.1 this reduces uniqueness of the stationary equilibrium to Light’s Theorem 1, that saving rises with the rate at every age, which is the subject of Section 7: it holds at Aiyagari’s calibration for relative risk aversion up to about 3.4, so the chain closes at his $\mu = 3$ and not at his $\mu = 5$.

6.4 Continuity, and existence

The last input is continuity of capital supply in the rate. In the infinite horizon this was a substantial development, because the stationary distribution itself had to be shown to move continuously. Here there is no stationary distribution: capital supply is a finite composition of finitely many policies, so all that is needed is that each policy moves continuously with the rate, and that continuity composes.

We take the route of the companion write-up and put the interest rate into the state, so that Berge’s theorem gives continuity in assets and the rate jointly rather than continuity in a parameter. That device was used there at the fixed point, but everything it rests on — that the feasible sets and the objectives agree on the slice — holds for an arbitrary continuation. So it works stage by stage, and here it is easier, because the slice of the augmented Bellman iterate is the iterate of the sliced programme by a plain induction, with no fixed point to identify (`sliceAt_bellman, sliceAt_augStageValue`). The augmented argmax at each stage is then a singleton, since the sliced one is by Carroll–Kimball concavity of the iterates (`argmax_aug_stage_eq_singleton`), so the stage policy is jointly continuous (`continuousOn_`

`augStagePolicy`) and agrees with the sliced one (`augStagePolicy_eq`). Running the cohort recursion in the augmented state (`augCohortAssets`) gives a composition of continuous maps that reads, at a rate in range, as capital supply at that rate (`augCohortAssets_eq, olgSupply_eq`).

Theorem 6.3 (Capital supply is continuous). *$r \mapsto K(r)$ is continuous on $[r_{lo}, r_{hi}]$.*

Formal statement: *`IncomeFluctuation.continuousOn_olgSupply`.*

Theorem 6.4 (Existence). *If capital supply lies below demand at the bottom of the interval and above it at the top, and demand is continuous, then the two meet: there is a rate in $[r_{lo}, r_{hi}]$ at which the capital the cohorts accumulate equals the capital the firm demands.*

Formal statement: *`IncomeFluctuation.exists_olgEquilibrium`.*

Both hypotheses of Theorem 6.4 are about the solved model. The rest of this section replaces them with inequalities that never mention the household's policy or value function.

6.5 Where the household can be

Proposition 6.9 inherits from Theorem 3.1 the hypothesis that the artefactual cap is slack, and the convenient way to state that is slack everywhere in $[0, \bar{a}]$. That hypothesis is not innocent, and it fails exactly where the floor is wanted.

The one-step bound gives $g_k(a, z) \leq (1 - \kappa_k)(Ra + y_{\max})$, so $[0, \bar{a}]$ is carried into itself only if $(1 - \kappa_k)(R\bar{a} + y_{\max}) \leq \bar{a}$, which needs $(1 - \kappa_k)R < 1$. As k grows κ_k falls to $1 - \mathfrak{p}/R$, so the requirement is $\mathfrak{p} < 1$, that is $\beta R < 1$: the household must be impatient. Above the rate of time preference no cap is forward invariant at any level, because a patient household sitting at the cap wants to save past it. The cap then binds, the Euler inequality under it is unavailable, and the slackness hypothesis is false rather than merely unverified. A theorem carrying it says nothing there.

This matters because a stationary life-cycle equilibrium sits above the rate of time preference. In the calibration of Section 6.10 with $\lambda = 1/\beta - 1 = 4.17\%$, capital supply crosses demand near 5%, and the rate at which the *provable* floor crosses demand is near 9%. Both are in the region where no cap is slack everywhere.

A cohort that begins life with nothing never visits the top of $[0, \bar{a}]$. The development therefore carries a family of regions, one per stage, that the stage policies map down through, and asks for slackness only on those.

Definition 6.5 (Reachable family). A reachable family assigns to each stage k a set of asset levels such that every set sits inside $[0, \bar{a}]$, the stage- $(k + 1)$ policy carries the stage- $(k + 1)$ set into the stage- k set, and $g_k < \bar{a}$ on every set.

Formal statement: *`IncomeFluctuation.StageRegions`.*

The whole interval $[0, \bar{a}]$ is a reachable family exactly when the cap is slack on it (`allRegions`), so nothing is lost: Theorem 3.1 and Proposition 6.9 are the constant-family instances of statements that hold on any reachable family (`stageConsumption_le_stageHumanWealth_on, le_stagePolicy_on`).

What a cohort visits is bounded by feasibility alone. Set $\rho_0 = 0$ and $\rho_{j+1} = R\rho_j + y_{\max}$: a household that cannot carry forward more than its cash on hand holds at most ρ_j after j periods of life, whatever the earnings draws.

Proposition 6.6 (A cohort's own regions). *If $\rho_{K+1} < \bar{a}$ then the sets $[0, \rho_{K-k}]$, empty for $k > K$, form a reachable family for a cohort of length $K + 1$ that starts life with nothing.*

Formal statement: *IncomeFluctuation.reachRegions.*

The bound is deliberately crude, using none of the damping that optimal saving supplies, and that is what keeps it free of the sandwich it is used to prove. Its cost is a larger cap. At $r = 10\%$ over sixty periods ρ_{K+1} is about 9,044 where the optimal path reaches about 82. Since the cap is an artefact of the state space rather than an economic object, a cap of ten thousand is as acceptable as a cap of one hundred, and unlike slackness everywhere it is a condition that can actually be met. The sandwich and the floor then hold along a cohort's path with no hypothesis on the cap beyond $\rho_{K+1} < \bar{a}$ (`stageConsumption_le_stageHumanWealth_reach`, `le_stagePolicy_reach`).

At the bottom of the rate interval the same crude bound does the other half of the work on its own. When $R < 1$ the sequence ρ_j converges to $y_{\max}/(1 - R)$, so cohort assets are bounded without any appeal to optimality: at $r = -7\%$ the bound is 42.1 against a capital demand of 56.3.

6.6 The propensity in closed form

Both bounds below are built from the finite-horizon propensity κ_k and from human wealth, and both need to know not just where κ_k sits but how fast it settles.

Proposition 6.7 (The finite-horizon propensity). *Write $q = p/R$. Then $\kappa_k \sum_{i \leq k} q^i = 1$, and hence $\kappa_k(1 - q^{k+1}) = 1 - q$: the propensity exceeds its infinite-horizon limit $1 - q$ by exactly $\kappa_k q^{k+1}$.*

Formal statement: *IncomeFluctuation.stageMPC_mul_mpcSum.*

, IncomeFluctuation.stageMPC_mul_one_sub_pow

The defining recursion $\kappa_{k+1} = \kappa_k R / (p + \kappa_k R)$ is exactly $1/\kappa_{k+1} = 1 + q/\kappa_k$, so the reciprocal accumulates the geometric series and nothing else. The same sum tames the slope that both bounds share. Writing $S_k = \sum_{i \leq k} q^i$, the recursion $G_0 = 1$, $G_{k+1} = 1 + G_k(1 - \kappa_{k+1})R$ has an age-dependent coefficient, but $W_k = G_k S_k$ obeys $W_{k+1} = S_{k+1} + pW_k$, which does not. Then $S_{k+1} \geq 1$ and $(1 - q)S_k \leq 1$ give the bound that matters.

Proposition 6.8 (A slope that grows). *$G_k \geq (1 - q) \sum_{j \leq k} b^j$.*

Formal statement: *IncomeFluctuation.mul_geomSum_le_cohortSlope.*

Where the crude $G_k \geq 1$ does not grow at all, this grows geometrically, and Section 6.11 reports what that is worth.

6.7 A floor on capital

Existence needs capital supply above demand at the top of the rate interval. In the infinite horizon the only floor available came from an ergodic identity that loses a factor $\|u'\|_\infty^2 / \text{Var}$. The life cycle needs nothing of the kind: the upper sandwich of Theorem 3.1 bounds consumption at every age, and what is not consumed is saved.

Proposition 6.9 (Affine floor on saving). *At every age and every state in the region,*

$$g_k(a, z) \geq (1 - \kappa_k)(y_z + Ra) - \kappa_k H_k(z).$$

Formal statement: `IncomeFluctuation.le_stagePolicy_on.`

`, IncomeFluctuation.le_stagePolicy`

Iterating it is not a matter of tracking a distribution over assets. An affine lower bound on a cohort's accumulated assets suffices, $F_k(z) + G_k a \leq A_k(a, z)$, with the slope G of Section 6.6 common to all earnings states and an intercept carrying one value per state:

$$F_0 = 0, \quad F_{k+1}(z) = \sum_{z'} \pi(z, z') F_k(z') + G_k [(1 - \kappa_{k+1}) y_z - \kappa_{k+1} H_{k+1}(z)].$$

Proposition 6.10 (Affine floor on accumulated assets). *On any reachable family, $F_k(z) + G_k a \leq A_k(a, z)$; hence capital per head is at least $\sum_z \nu_z F_K(z) / (K + 1)$.*

Formal statement: `IncomeFluctuation.cohortFloor_add_mul_le_cohortAssets.`

`, IncomeFluctuation.le_olgCapital`

The state dependence of F is the whole content, not a refinement. Reading Proposition 6.9 at the worst earnings state instead, which is the only way to get a bound uniform in z , gives a floor that is identically zero at every rate in the calibration of Section 6.10, against 4.75 at $r = 10\%$ for the state-dependent one.

Averaging against weights the chain leaves alone collapses the vector recursion to a scalar one, and the sign of each increment then has an exact answer.

Theorem 6.11 (The floor rises exactly when the household is patient). *Let $\bar{y}$ and $\bar{H}_k$ be the averages of earnings and of human wealth under weights ν satisfying $\nu\pi = \nu$. If $\bar{p} \geq 1$, equivalently $\beta R \geq 1$, then $\kappa_k \bar{H}_k \leq (1 - \kappa_k) \bar{y}$ at every age, so the aggregate floor never falls and is never negative.*

Formal statement: `IncomeFluctuation.stageMPC_mul_sum_stageHumanWealth_le.`

`, IncomeFluctuation.sum_mul_cohortFloor_mono`

The increment of the aggregate floor is $G_k \kappa_{k+1} \bar{y}$ times the slack in that inequality, and the inequality reverses below $\bar{p} = 1$. So the provable floor is flat under the rate of time preference and climbs above it, which is why the upper endpoint of the rate interval has to be taken above $\lambda = 1/\beta - 1$. The proof is one induction: writing $v_k = (1 - \kappa_k)/\kappa_k$ and $u_k = \bar{H}_k/\bar{y}$, both satisfy first-order recursions, $v_{k+1} = (\bar{p}/R)(1 + v_k)$ and $u_{k+1} = (1 + u_k)/R$, from the same starting point, and $\bar{p} \geq 1$ keeps the first above the second.

One thing the affine form cannot do is refuse to claim negative assets. Its intercept goes badly negative: at the calibration, 160 of 413 entries are below zero and the worst is -62 ,

with F_K running from -58.8 at the lowest earnings to $+685$ at the highest. A cohort's accumulated assets cannot be negative, so all of that is slack the aggregate inherits. Reading the floor as a function of assets rather than as an affine form fixes it, because the saving floor can then be clamped at zero where it belongs:

$$\Lambda_0(z, a) = a, \quad \Lambda_{k+1}(z, a) = a + \sum_{z'} \pi(z, z') \Lambda_k(z', \Phi_{k+1}(z, a)),$$

with Φ the clamped saving floor of Section 7.13. The proof needs only that Λ_k is increasing in assets, which the clamp preserves.

Proposition 6.12 (The clamped floor). *$\Lambda_k(z, a) \leq A_k(a, z)$ on any reachable family, so capital per head is at least $\sum_z \nu_z \Lambda_K(z, 0)/(K + 1)$, which is at least what Proposition 6.10 gives.*

Formal statement: *IncomeFluctuation.cohortFloorAt_le_cohortAssets.*
, IncomeFluctuation.le_olgCapital_at

6.8 A ceiling on capital

Existence also needs capital supply below demand at the bottom of the interval, and there three bounds are available, each asking for more and giving more.

The cheapest uses the budget and nothing else. A household cannot carry forward more than its cash on hand, so an asset level B with $RB + y_{\max} \leq B$ is never left once entered.

Proposition 6.13 (A forward-invariant level caps supply). *If $0 \leq B \leq \bar{a}$ and $RB + y_{\max} \leq B$, then capital per head is at most B .*

Formal statement: *IncomeFluctuation.olgCapital_le_of_invariant.*

The second says the household will not want to carry that much. Consuming at least $\kappa_k(m + H_k(z))$ by the constraint-incidence bound of Theorem 7.8, it saves at most $(1 - \kappa_k)(y_z + Ra) - \kappa_k H_k(z)$.

Proposition 6.14 (The incidence ceiling). *If $B \leq \bar{a}$ and, at every age and earnings state, $(1 - \kappa_j)(y_w + RB) - \kappa_j H_j(w) \leq B$, then capital per head is at most B .*

Formal statement: *IncomeFluctuation.olgCapital_le_of_incidence.*

Both bound *every* age's assets by one level, and so pay $(K + 1)B$ for a cohort that holds nothing when it is born, nothing when it dies, and anything like B only in between. The floor never made that mistake, and the third ceiling does not either: it is the same affine object, built from the same slope G , with the saving ceiling in place of the saving floor,

$$\bar{F}_0 = 0, \quad \bar{F}_{k+1}(z) = \sum_{z'} \pi(z, z') \bar{F}_k(z') + G_k[(1 - \kappa_{k+1})y_z - \kappa_{k+1} H_{k+1}(z)],$$

with H here the capped, risk-adjusted human wealth of Theorem 7.8.

Proposition 6.15 (The affine ceiling). *$A_k(a, z) \leq \bar{F}_k(z) + G_k a$, and hence capital per head is at most $\sum_z \nu_z \bar{F}_K(z)/(K + 1)$.*

Formal statement: *IncomeFluctuation.cohortAssets_le_cohortCeil.*
, IncomeFluctuation.olgCapital_le_cohortCeil

Replacing a worst case over ages by an average over them is worth a factor of fourteen, and it keeps working above the rate of time preference, where no forward-invariant level exists at all.

	$r = -7\%$	$r = -4\%$	$r = 0\%$	$r = \lambda$	$r = 5\%$
Capital demand	56.25	14.06	7.03	4.62	4.33
Proposition 6.13, the budget	42.52	74.41	—	—	—
Proposition 6.14, incidence	22.17	35.11	—	—	—
Proposition 6.15, affine	1.54	2.28	3.86	5.98	6.47
True capital supply	0.64	1.03	1.98	3.78	4.25

6.9 Existence from primitives

The two ends can now be put together, and the statement has taken four forms as the bounds improved. All four are in the development; the last is the one to read.

Theorem 6.16 (Existence from primitives). *Fix a rate interval and continuous demand D . Suppose there is a B with $0 \leq B \leq \bar{a}$, $(1 + r_{lo})B + y_{\max} \leq B$ and $B \leq D(r_{lo})$; and suppose $\rho_{K+1} < \bar{a}$ and $D(r_{hi}) \leq \sum_z \nu_z F_K(z)/(K+1)$ at r_{hi} . Then there is a rate in the interval at which the capital the cohorts accumulate equals the capital the firm demands.*

Formal statement: *IncomeFluctuation.exists_olgEquilibrium_rate.*

, *IncomeFluctuation.exists_olgEquilibrium_of_bounds*

Nothing in those hypotheses refers to the household's decisions. They are inequalities among the discount factor, the interest rate, the earnings process, the horizon and the demand curve, with ρ , F and G explicit finite recursions in those primitives. Three refinements follow. Proposition 6.8 turns the floor into a closed form with no recursion left to run (*IncomeFluctuation.exists_olgEquilibrium_of_geomBounds*, Theorem 6.19 below); and Proposition 6.14 replaces the budget ceiling (*IncomeFluctuation.exists_olgEquilibrium_of_sharpBounds*). The affine ceiling replaces it again, and at that point no auxiliary B and no invariance condition survive.

Theorem 6.17 (Existence, both ends read the same way). *With the ceiling at r_{lo} and the floor at r_{hi} both taken as weighted averages of affine bounds on a cohort's accumulated assets, a stationary equilibrium exists.*

Formal statement: *IncomeFluctuation.exists_olgEquilibrium_affine.*

Reading the floor at the assets held as well gives the tightest form.

Theorem 6.18 (Existence with both ends read at the assets held). *With the affine ceiling at r_{lo} and the clamped floor at r_{hi} , a stationary equilibrium exists.*

Formal statement: *IncomeFluctuation.exists_olgEquilibrium_clamped.*

The two boundary conditions are then the same object built twice, once from the saving ceiling and once from the saving floor, sharing a slope. Each half asks for what it can have: at the bottom the artefactual cap is slack, which holds below the rate of time preference; at

the top it cannot be, and the reachable family of Section 6.5 stands in. So the low endpoint must sit below λ and the high endpoint above it, and Section 6.11 shows a bracket that does.

For completeness, the closed-form variant.

Theorem 6.19 (Existence in closed form). *Under the conditions of Theorem 6.16, with the floor hypothesis replaced by the closed-form expression of Proposition 6.8 and with $1 \leq p \leq R$ at r_{hi} , a stationary equilibrium exists.*

Formal statement: `IncomeFluctuation.exists_olgEquilibrium_of_geomBounds`.

6.10 What the numbers say

The calibration is the one used throughout: $J = 60$, $\beta = 0.96$, relative risk aversion 3, the Tauchen discretisation of Aiyagari’s earnings process with $\rho = 0.9$ and $\sigma = 0.4$, and a Cobb–Douglas technology with $\alpha = 9/25$ and $\delta = 2/25$, so capital demand is $\alpha/((1-\alpha)(r+\delta))$ and $\lambda = 4.17\%$. Capital supply is the average asset holding across the sixty cohorts, each born with nothing. The last row is not the solved supply but the floor of Proposition 6.10 (`numerics/olgexist.m`).

r	−4%	0%	λ	6%	8%	10%	15%
Capital demand	14.06	7.03	4.62	4.02	3.52	3.13	2.45
Capital supply	1.03	1.98	3.78	4.85	6.18	7.69	12.49
Provable floor	0.00	0.00	0.00	1.52	3.13	4.75	9.40

Supply crosses demand at 5.09%, above the rate of time preference, which is what Section 6.5 has to accommodate; and the floor is flat below λ and climbs above it, which is Theorem 6.11. The optimal path reaches assets of 82 at $r = 10\%$, against the crude reachable bound of 9,033; the constants are ordinary, unlike the stationary case, where the corresponding instantiation put r_{hi} within 10^{-27} of λ .

6.11 The theorem checked, hypothesis by hypothesis

Take $r_{lo} = 2\%$, $r_{hi} = 8\%$ and an artefactual cap of 10^4 . Every hypothesis of Theorem 6.18 holds (`numerics/clampcheck.m`): the cap is forward invariant at r_{lo} with margin 134; the ceiling is 4.81 against a demand of 5.63; a lifetime of feasible saving reaches 3,730 against the cap; and the floor is 3.80 against a demand of 3.52. The bracket contains the true equilibrium at 5.09%.

Sweeping the endpoints shows what each bound is worth. The low endpoint may be taken as high as

Ceiling used	highest r_{lo}	Floor used	lowest r_{hi}
Proposition 6.13, the budget	−6.73%	closed form, crude $G_k \geq 1$	610%
Proposition 6.14, incidence	−5.85%	closed form, Proposition 6.8	18.2%
Proposition 6.15, affine	+2.78%	Proposition 6.10, affine	8.38%
		Proposition 6.12, clamped	7.67%

So the tightest bracket the development gives is [2.78%, 7.67%], and the absolute limit is 5.09% from both sides, since no bound on capital supply can cross demand before supply does. The closed-form floor is reported because it needs no recursion at all; at a calibrated rate it is much the weaker, giving 0.48 at $r = 10\%$ against a true floor of 4.75.

6.12 Where the bounds stop

Both ends are now within a few points of the limit, and neither will move further with the tools here. The reasons are worth recording, because they are the same reason twice.

At the top, the floor's entire loss is in one inequality. Running the floor's own recursion on the cohort's *true* mean consumption path reproduces aggregate capital exactly, 6.180 against 6.180 at $r = 8\%$ (`numerics/diag.m`), so none of it is in the aggregation, the affine bound or the earnings weights: all of it is in $c_k \leq \kappa_k(m + H_k(z))$. That inequality is exact at the end of life and a fifth loose at the start, and the leverage is what hurts. At birth the household earns 1 and consumes 0.698, so it saves 0.302; the bound allows 0.858 and leaves a provable saving of 0.142. Overstating consumption by a fifth halves the saving.

Three ways of tightening it fail. Truncating the horizon gives $c_k \leq \kappa_j(m + H_j(z))$ for every $j \leq k$, all valid, but $\kappa_j(m + H_j)$ is decreasing in j so the minimum is the full-horizon branch already in use. Substituting the power mean that the Euler step actually produces, in place of the arithmetic mean the proof relaxes it to, is not legitimate: the step shifts by gross saving, and a power mean with negative exponent is superadditive, so the shift comes out larger inside than outside, which is the wrong way for an upper bound; carried through anyway it violates the true consumption function by 0.30 and gives a floor of 7.45 against a true supply of 6.18. Putting the shift inside the mean avoids superadditivity but forces tomorrow's bound to be read at a saving ceiling, decoupling the budget identity from the Euler step, and loses more than the power mean gains, 0.956 against 0.858 at birth. That last one also measures the sensitivity: with the *true* saving the same recursion reproduces the truth exactly, and the ceiling it must use instead is only 1.19 times the truth at birth.

6.13 And the same obstruction from the other side

At the bottom there is no structural slack left. Reading the ceiling at the assets held, with the same human wealth, gives the identical number at every rate, 4.8147 against 4.8147 at $r = 2\%$ (`numerics/ceilat.m`): given its one-step bound the affine ceiling is already exact. What remains is the one-step bound, and there the asset-dependent incidence bound of Section 7.13 does help, because the affine ceiling averages over states and so picks up the low earnings states where the cap bites. It is worth 0.7 points, the same as the clamp gave at the top: the low endpoint would move from 2.78% to 3.47%.

The proof does not follow. Taking those points needs the ceiling read at the assets held, which needs its argument, the saving ceiling, to increase with assets. It does, with not one decreasing step over the grid, and the truth sits exactly on the bound the induction would need: the largest slope of the asset-dependent human wealth is 26.147 and the cap's slope is 26.147, because it is the cap branch that binds. But carrying that bound through the

recursion means passing a perturbation through the power mean, and a power mean with negative exponent is not 1-Lipschitz in the sup norm. It is superadditive, so a uniform shift comes out amplified, here by 9.1; over fifty-nine ages the natural induction would allow the slope to grow by 4×10^{56} , against a truth that does not grow at all (`numerics/slopes.m`). The cheap way round, reading the asset-dependent human wealth at zero assets, is valid and monotone by construction and moves the endpoint from 2.78% to 2.81%.

So superadditivity of the power mean blocks both ends: at the top it stops the mean being substituted into the consumption upper bound, and at the bottom it stops a slope bound being propagated past it. It is the same obstruction that closed the rate-response bound in the stationary companion.

7 Theorem 1 in the life cycle: how far it goes

The programme of Section 11 rests on Light’s Theorem 1, that saving rises with the interest rate at every state. In the infinite horizon this is open above logarithmic utility. The life cycle was the reason for hope: a finite backward induction from an explicit terminal condition, with the per-step loss compounding only J times instead of to a fixed point. This section reports what the numbers say, at $\beta = 0.96$, a seven-state Tauchen chain with $\rho = 0.9$ and $\sigma = 0.4$, a zero borrowing limit, $J = 60$, and a flat earnings profile (`numerics/lc_check.m`, `lc_robust.m`, `lc_induct.m`, `lc_verify.m`, `lc_bound.m`, `lc_crit.m`).

7.1 The statement is its own terminal condition

The budget constraint makes two statements identical. Since $a' = Ra + y(z) - c$ at either rate,

$$g_k(a, z; r_2) \geq g_k(a, z; r_1) \iff c_k(a, z; r_2) - c_k(a, z; r_1) \leq (r_2 - r_1)a,$$

so Theorem 1 *is* the affine bound on the rate response of consumption with slope one and zero intercept. That bound holds with equality in the last period of life, where $c_0 = y(z) + Ra$. The question is therefore whether the terminal condition propagates backward, and it is a question about one inequality rather than about a pair of constants to be chosen.

7.2 It propagates, up to a critical risk aversion

At $\mu = 1$ and $\mu = 3$ it propagates exactly. Over all sixty stages and the whole wealth region reachable from zero wealth at birth, the largest violation of $g_k^{(1)} \leq g_k^{(2)}$ is zero, and the consumption bound holds to five parts in 10^{14} ; refining the asset grid from 2000 to 16000 points changes nothing. At $\mu = 5$ it fails. The failure is a genuine feature of the household’s problem, not of the solution method: the violation is identical at every grid size, and endogenous-grid and value-function iteration with a pure grid search agree on it to the resolution of the grid. Saving at zero wealth in the highest income state, in the first period of a sixty-period life, falls as the rate rises:

r	2%	3%	4%	5%	6%
$g(0, z_{\max})$, endogenous grid	1.7783	1.7638	1.7452	1.7234	1.6993
$g(0, z_{\max})$, value iteration	1.7806	1.7606	1.7406	1.7206	1.7006

Between the two lies a critical degree of risk aversion, and it is sharp. Bisecting on μ for the largest violation over all stages gives:

Horizon J	$r = 2\%$	$r = 4\%$	$r = \lambda = 4.17\%$
20	3.79	3.53	3.51
40	3.77	3.39	3.36
60	3.77	3.39	3.36

Theorem 1 holds throughout the life cycle below the critical value and fails above it. The value is stable in the horizon beyond about forty periods, and falls as the rate rises towards the rate of time preference. Aiyagari's $\mu = 3$ is inside it with little room, at $\mu = 3.39$ against 3; his $\mu = 5$ is well outside. So the life cycle buys one of the two open cases and not the other, and it buys $\mu = 3$ only at rates below the rate of time preference, which is where the equilibrium lies.

7.3 Why the natural induction does not close

The backward induction that the transfer lemma suggests compares the two Euler equations at a state where the lower-rate household saves. Writing $b = g_{k+1}^{(1)}(a, z)$ for the pivot, the induction hypothesis and monotonicity of c_k give $c_k^{(2)}(b') \leq c_k^{(1)}(b) + \Delta r b$ at the relevant points, and the step closes provided

$$R_2 \mathbb{E}[(c_k^{(1)}(b) + \Delta r b)^{-\mu}] \geq R_1 \mathbb{E}[c_k^{(1)}(b)^{-\mu}] \left(1 + \frac{\Delta r a}{c_{k+1}^{(1)}(a, z)}\right)^{-\mu}.$$

To first order this asks $\mu(b/c_k^{(1)}(b) - a/c_{k+1}^{(1)}(a, z)) \leq 1/R$: the ratio of assets to consumption may rise from one age to the next by at most $1/(\mu R)$. It does not. At zero wealth in the highest income state the left-hand side is large, because such a household saves a great deal against a certain fall in income, while the right-hand side is about one. The margin is negative from the third stage of life onward for every degree of risk aversion we tried, logarithmic utility included, at -0.12 for $\mu = 1$, -2.86 for $\mu = 3$ and -6.29 for $\mu = 5$ per unit of the rate gap. Keeping the term the step discards, the marginal propensity to consume times the size of the hypothetical violation, bounds the violation by a multiple of Δr rather than eliminating it, which is not enough: the aggregate difference the equilibrium argument must sign is itself of order Δr .

7.4 The reduction, formalised, and its verdict

The backward induction above is proved in Lean, in the form the life cycle makes natural. Because the budget constraint makes the two statements identical (`stagePolicy_le_iff_`

consumption_{le}), the terminal stage needs no hypothesis at all: in the last period both households consume their cash on hand and save nothing (stagePolicy_mono_zero). Corners need none either: if the lower-rate household saves nothing there is nothing to violate. What the step needs is one inequality per age and state, which we call the *age-indexed slack condition*: writing b for the lower-rate household's saving at that state and $C'(z')$ for its consumption next period,

$$R_1 \sum_{z'} \pi(z, z') C'(z')^{-\mu} \leq R_2 \sum_{z'} \pi(z, z') (C'(z') + \Delta r b)^{-\mu} \left(1 + \frac{\Delta r a}{c}\right)^\mu \quad (1)$$

(StageSlack).

Theorem 7.1 (The reduction). *If (1) holds at every age and every state in the region, then saving rises with the rate at every age, and aggregate capital rises with the rate.*

Formal statement: *IncomeFluctuation.stagePolicy_mono_step,*

IncomeFluctuation.stagePolicy_mono_of_stageSlack, IncomeFluctuation.olgCapital_le_of_stageSlack.

With Theorem 6.2 this closes the chain: (1) at every age gives a unique equilibrium rate. The question is therefore whether (1) is true, and the answer is the one reported above, now stated sharply. Condition (1) is exactly the margin tabulated in the previous subsection. At $\mu = 3$ it holds for the last three or four periods of life, with margins +0.68, +0.42 and +0.18 per unit of the rate gap, and fails from about the fifth-to-last, at -0.23 , deteriorating to -2.86 at birth. At $\mu = 1$ it holds for most of the life and still fails for the youngest ages, at -0.12 at birth.

That last observation is the important one. The condition fails even at logarithmic utility, where Theorem 1 is certainly true. So the failure is not a symptom of high risk aversion: the pointwise Euler comparison is simply the wrong instrument in a life cycle, and no sharpening of it that we could find repairs the young ages. Keeping the term the step discards, the marginal propensity to consume times the size of the hypothetical violation, only bounds the violation by a multiple of Δr , and both refinements we tried — an upper bound on the propensity to consume from concavity and the sandwich, and a self-improving bootstrap on the size of the violation — converge to a positive bound rather than to zero.

So $\mu = 3$ is not delivered by this route, and $\mu = 1$ is not either, although the companion write-up proves the latter by the rescaling argument, which is an induction on Bellman iterates and so carries over to the life cycle unchanged. What the reduction does deliver is the end of life, and the exact inequality that any proof must beat.

So the state of the programme is this. Theorem 1 is true in the life cycle for $\mu = 3$, and the statement to prove has the cleanest possible form, its own terminal condition propagated; the reduction to one inequality per age is proved; and that inequality is false for the young, at every degree of risk aversion. What is missing is an argument at zero wealth that does not go through a pointwise comparison of the two Euler equations. The same zero-wealth state is where the infinite-horizon slack condition was tightest, which suggests the two problems share a single missing ingredient rather than two.

7.5 It is not precautionary saving

A household at zero wealth in the highest earnings state saves heavily, and the obvious diagnosis is precautionary: it is insuring against the fall in earnings it expects. If that were the mechanism, precautionary-saving theory would be the place to look for the missing argument. It is not the mechanism. The same reversal happens in the *certainty-equivalent* household, which faces the conditional mean earnings path and no risk at all, and there everything is in closed form (`numerics/precaut2.m`, `durcheck.m`).

d(saving)/dr at zero wealth	$\mu = 1$	$\mu = 3$	$\mu = 5$	$\mu = 8$
Certainty equivalent (mean path, no risk)	+21.3	+2.81	-0.85	-2.91
With risk	+17.1	+0.92	-2.05	-3.43

The sign flips between $\mu = 3$ and $\mu = 5$ in both columns. Bisecting, the certainty-equivalent threshold is $\mu = 4.33$ against $\mu = 3.39$ with risk: risk moves the threshold by about one unit of risk aversion and does not create the effect.

Without risk the mechanism is exact. A household with J periods, no borrowing constraint and CRRA utility has $c_{t+1} = (\beta R)^{1/\mu} c_t$ and a lifetime budget, so

$$c_0 = \frac{W(R)}{D(R)}, \quad W(R) = \sum_t y_t R^{-t}, \quad D(R) = \sum_t \beta^{t/\mu} R^{t/\mu - t}$$

(`consumption_mul_pathPrice`), and saving at zero wealth rises with the rate exactly when $W(R_2)D(R_1) \leq W(R_1)D(R_2)$ (`consumption_antitone_iff`). Both terms fall with the rate, W because future earnings are discounted harder and D because the optimal path gets cheaper, and the comparison is between the two elasticities. Differentiating,

$$\frac{d \log c_0}{d \log R} = -T_y + \left(1 - \frac{1}{\mu}\right) T_c,$$

with T_y and T_c the Macaulay durations of the two streams, so c_0 falls, and saving rises, exactly when $T_y > (1 - 1/\mu)T_c$. The predicted elasticity matches the exact one to three decimals at every μ we checked. At Aiyagari's numbers $T_y = 14.31$ and $T_c \approx 18.6$, nearly independent of μ , so the condition is $1 - 1/\mu < 0.769$, that is $\mu < 4.33$.

The economics is then plain, and has nothing to do with risk. A higher interest rate discounts future earnings more heavily, which makes the household poorer and pushes consumption down; it also makes a rising consumption path cheaper, which pushes consumption today down only to the extent that the household is willing to tilt its path. A more risk averse household will not tilt, so the second force weakens as μ grows while the first does not, and past $\mu \approx 4$ the household responds to a higher rate by consuming more today. The income stream simply has to be longer than the consumption weights, discounted by the strength of the tilt.

Theorem 7.2 (No risk, and risk aversion at most one). *If $\mu \leq 1$ the comparison holds term by term, whatever the earnings profile: the certainty-equivalent household always saves more at a higher rate.*

Formal statement: *CertaintyEquivalent.humanWealth_mul_pathPrice_le_of_le_one*,
CertaintyEquivalent.consumption_antitone_of_le_one.

For $\mu > 1$ it is the duration comparison, which is a finite computation in the primitives. So a precautionary argument near zero wealth would be aimed at the wrong target: what has to be controlled is the elasticity of human wealth against the elasticity of the price of the consumption path, and risk enters only as a correction to it.

7.6 The criterion in elasticity units, and the size of the risk correction

The certainty-equivalent criterion is cleanest when written in the intertemporal elasticity of substitution rather than in risk aversion. Since it is linear in $1/\mu$,

$$T_y > \left(1 - \frac{1}{\mu}\right)T_c \quad \Longleftrightarrow \quad \frac{1}{\mu} > 1 - \frac{T_y}{T_c},$$

so the elasticity of substitution must exceed one minus the ratio of the two durations. At Aiyagari's numbers that right-hand side is 0.2306, giving $\mu < 4.33$ as before, and the identity holds to four decimals against the bisected threshold.

With risk the threshold moves down, and the move is what has to be bounded. Write

$$\Delta := \frac{1}{\mu^{\text{risky}}} - \frac{1}{\mu^{\text{CE}}}$$

for the gap in elasticity units and $s = \sigma\sqrt{1 - \rho^2}$ for the conditional standard deviation of log earnings, the dispersion the household actually faces one period ahead. Measuring both thresholds at the same state over fifteen calibrations, $\rho \in \{0, 0.6, 0.9\}$ and σ from 0.025 to 0.4 (numerics/gapscale.m, gapfit.m):

$\rho = 0.9, \sigma$	0.025	0.05	0.1	0.2	0.4	fitted exponent
Δ	0.0034	0.0070	0.0144	0.0300	0.0612	
Δ/s	0.315	0.320	0.330	0.344	0.351	1.04

The gap is first order in the conditional standard deviation, not second order: fitting $\Delta \sim s^b$ gives $b = 1.04$ at $\rho = 0.9$, 0.89 at $\rho = 0.6$ and 0.93 at $\rho = 0$, and the ratio Δ/s stays between 0.196 and 0.351 across every calibration we ran. That first-order behaviour is itself informative: a precautionary effect on saving would be second order in the dispersion, whereas the probability that the borrowing constraint binds in a low-earnings state next period moves at first order, and that is our reading of where the correction comes from.

Taking the largest ratio observed as the bound, $\Delta \leq 0.36 s$, gives a criterion entirely in the primitives:

$$\frac{1}{\mu} \geq 1 - \frac{T_y}{T_c} + 0.36 \sigma \sqrt{1 - \rho^2}. \quad (2)$$

At Aiyagari's calibration the right-hand side is $0.2306 + 0.0628 = 0.2934$, that is $\mu \leq 3.41$, against the 3.39 that bisection on the full life-cycle model gives. The bound therefore reproduces the measured threshold to within two per cent, and it says the same thing the earlier sections said: the calibration he used for $\mu = 3$ is inside, and $\mu = 5$ is not.

The certainty-equivalent half, for every risk aversion

The first inequality of (2) was a theorem only for $\mu \leq 1$, where the exponent $\theta = 1 - 1/\mu$ is nonpositive and every factor moves the same way. It is now a theorem for every μ . Writing $\lambda = R_1/R_2 < 1$, the criterion $W(R_2)D(R_1) \leq W(R_1)D(R_2)$ compares how the same weights respond to λ^t against $(\lambda^t)^\theta$. Each λ^t lies in $[\lambda^{J-1}, 1]$, and on that interval the concave $x \mapsto x^\theta$ lies above its chord, which is affine; replacing the power by its chord leaves a comparison of two sums linear in λ^t .

Theorem 7.3 (The chord criterion). *For $\mu > 1$, if $W(R_2)D(R_1) \leq W(R_1)(\iota D(R_1) + \varsigma D_\lambda(R_1))$, where ι and ς are the intercept and slope of that chord and D_λ is the path price with each term discounted once more by λ^t , then initial consumption falls with the rate.*

Formal statement: *CertaintyEquivalent.consumption_antitone_of_chord.*

, CertaintyEquivalent.chord_le_pathPrice

The replacement costs almost nothing, because it is first-order exact: as $R_2 \rightarrow R_1$ it reproduces $T_y \geq (1 - 1/\mu)T_c$ exactly. At Aiyagari's numbers the chord threshold is 4.3263 against the exact 4.3267, and the cost grows only with the size of the rate step, to 0.04 at $\Delta r = 10^{-3}$ and 0.40 at 10^{-2} (`numerics/chord.m`). Together with the $\mu \leq 1$ theorem this settles the certainty-equivalent household at every risk aversion.

The risk half, and why it resists

Two caveats remain. The constant 0.36 is measured over fifteen calibrations and is not proved. And both are statements at zero wealth in the highest earnings state, which is where the life-cycle model binds; the reduction of Theorem 7.1 is what would carry them to the rest of the state space, and its own condition fails, so this is a sharp description of where the truth lies rather than a proof of it.

It is worth being precise about why the constant resists. What has to be bounded is a difference of *rate responses* between two problems, the risky one and its certainty equivalent, and the threshold is the rate at which one of those responses changes sign. No bound on the *level* of consumption helps with that, however sharp: at Aiyagari's calibration the sharpest sandwich available has width 0.30 to 0.48 at the state in question, while the response itself runs from -2.48 at $r = 2\%$ through zero to $+0.29$ at $r = 6\%$ (`numerics/scale.m`). A bound on where consumption lies says nothing about how it moves. What is needed is a bound on the rate response itself, and that is exactly the object the companion write-up proved cannot be had from the affine induction, whose per-step multiplier is at least $1/p > 1$. The two halves of (2) therefore stand on different ground: the first is now a theorem for every risk aversion, and the second is blocked by the same obstruction that closed the rate-response bound.

7.7 The channel is the borrowing constraint

The first-order scaling points at the constraint rather than at precaution, and the point can be tested directly: loosen the borrowing limit and watch the threshold. If the constraint is

the channel, the risky threshold should walk back to the certainty-equivalent one as the limit approaches the natural one, $y_{\min}/r$. At Aiyagari’s calibration it does (`numerics/borrow.m`), and it does so almost linearly:

Borrowing limit	0	0.5	1	2	3	4	5	6
Threshold μ	3.42	3.50	3.58	3.74	3.89	4.05	4.21	4.38
Share of the gap closed	0%	11%	21%	40%	57%	74%	90%	104%

The natural limit here is $y_{\min}/r = 6.75$, and by a limit of 6 the gap is gone: the risky household’s threshold, 4.38, is the certainty-equivalent household’s, 4.33. So at this calibration the borrowing constraint accounts for the whole of the correction, and precaution accounts for none of it. Precaution is certainly present — at the loose limit the risky household still saves 1.62 against the certainty-equivalent 1.46, a precautionary excess of eleven per cent — but it moves the level of saving without moving the sign of its response to the interest rate. Across the other calibrations we ran (`numerics/borrow3.m`) loosening the limit always closes the gap and sometimes overshoots it, so the constraint is the dominant channel rather than an exact decomposition: with ample borrowing, risk can make Theorem 1 easier than certainty equivalence rather than harder.

For the programme this is the useful part. The missing ingredient for $\mu = 3$ is not a precautionary bound but a bound on what the borrowing constraint does to the comparative static, and that is a more tractable object: without a binding constraint the household’s Euler equation holds with equality at every state and Theorem 1 at zero wealth reduces to the duration criterion, which is proved outright for $\mu \leq 1$ and is an explicit finite computation in the primitives otherwise. What remains is the constrained states, which at this calibration are about eight per cent of the state space at every age.

7.8 The constraint channel is monotone, and needs no condition

The comparative static in the borrowing limit itself is a theorem, and unlike the one in the interest rate it needs nothing beyond the standing assumptions.

Theorem 7.4 (A tighter limit means less consumption and more saving). *Two households share every primitive and differ only in the borrowing limit, $f_2 \leq f_1$. Then at every age and every state in the tighter household’s region it consumes less and saves more.*

Formal statement: `IncomeFluctuation.finiteConsumption_le_of_floor_le`,
`IncomeFluctuation.le_finitePolicy_of_floor_le`.

The induction is the one of Theorem 7.1 with the interest rate held fixed, and that is exactly what removes the difficulty. Suppose the claim holds with k periods to come and fails at $k + 1$, so the tighter household consumes strictly more and therefore saves strictly less. Its Euler inequality under the cap, the looser household’s Euler inequality above its floor — available because the looser household saves strictly more than the tighter one and hence strictly more than either floor — the induction hypothesis, and monotonicity of consumption in wealth chain into $u'(c_1) \geq u'(c_2)$, contradicting $c_1 > c_2$. The gross return appears on both

sides and cancels, so there is no substitution term to control and no analogue of (1) is needed. Tightening the limit raises the marginal value of saving unambiguously, and the terminal stage is again free: a household with a lower floor runs down further in its last period.

That is the one-sided half of what the numbers showed, and it says the constraint channel is better behaved than the rate channel rather than merely different. What it does not yet give is a *size*: Theorem 7.4 orders the two consumption functions without bounding the distance between them, and it is the distance that turns into the 0.9 units of risk aversion between the constrained and certainty-equivalent thresholds. The natural next object is therefore a modulus for Theorem 7.4, and it has a clean form.

Theorem 7.5 (The modulus: a relaxation of Δ is worth at most Δ of wealth). *Let $\Delta = f_1 - f_2 \geq 0$ and $r \geq 0$. Then at every age and every state,*

$$c_k^{f_1}(x) \leq c_k^{f_2}(x) \leq c_k^{f_1}(x + \Delta).$$

Formal statement: *IncomeFluctuation.finiteConsumption_le_of_floor_le, IncomeFluctuation.finiteConsumption_le_shift.*

The household that may borrow down to f_2 therefore behaves, at wealth x , no more freely than the household confined to f_1 would at wealth $x + \Delta$: the entire value of the extra borrowing capacity is bounded by the capacity itself. The proof is the same backward induction, carried out at shifted arguments, and the shift pays for itself twice. The terminal stage holds with slack, since the looser household gains Δ by running down further while the shifted comparison gives it $(1 + r)\Delta$; and the interiority the step needs comes free, because the failure hypothesis forces the tighter household's saving to exceed the looser one's by more than Δ , hence to exceed $f_2 + \Delta = f_1$.

Two consequences. Since saving never falls with wealth, the sandwich gives the plain bound $0 \leq c_k^{f_2}(x) - c_k^{f_1}(x) \leq (1 + r)\Delta$, so the effect of the limit on consumption is at most the relaxation, up to one period's interest. And because the bound is stated as a shift in wealth rather than as an additive constant, it composes with the minimal propensity to consume: whatever upper bound on the propensity to consume out of wealth one has at age k , the modulus turns it directly into a bound on the consumption gap, and it is that product which becomes the gap between the constrained and certainty-equivalent thresholds.

7.9 The upper bound on the propensity to consume, and what it is worth

The missing piece is an upper bound on the propensity to consume out of wealth. The repository has the minimal propensity, which is a lower bound; the upper one follows by the same two-point Euler comparison, with no derivative taken. Write the bound as a secant slope in assets, so that s_k bounds $c_k(a_2) - c_k(a_1)$ by $s_k(a_2 - a_1)$.

Theorem 7.6 (The propensity to consume falls with the horizon). *In the last period the slope is exactly R (stageConsumption_sub_zero). If at stage k the slope is at most s and consumption is at least $c_{\min}$, then at a state where the poorer household saves, the slope at*

stage $k + 1$ is at most

$$R \cdot \frac{\nu}{1 + \nu}, \quad \nu = \frac{cs}{c_{\min}},$$

with c the poorer household's consumption.

Formal statement: *IncomeFluctuation.stageConsumption_sub_le_of_slope*.

The step is the two Euler inequalities at the two wealth levels: the induction hypothesis bounds the rise in tomorrow's consumption by $s\delta$ with δ the difference in saving, which bounds the fall in marginal utility by the factor $(1 + s\delta/c_{\min})^{-\mu}$, and the budget identity $c_2 - c_1 = R(a_2 - a_1) - \delta$ turns that into the stated slope. Since $\nu \mapsto \nu/(1 + \nu)$ is a contraction towards $1 - 1/\nu$, the slope falls from R towards $R(1 - c_{\min}/(cs))$ as the horizon lengthens: the longer the life, the less of an extra unit of wealth is spent at once, which is the qualitative fact one wants.

It is not sharp enough. At Aiyagari's calibration, read at zero wealth in the highest earnings state, consumption is about 1.32 and the worst consumption next period about 0.88, so $\nu \approx 1.5s$ and the recursion settles at a propensity of about 0.36, against a true propensity there of about 0.04. Fed through Theorem 7.5 it bounds the effect of the borrowing limit on consumption by about 2.5 where the truth is about 0.04, a factor of sixty. The loss is the ratio of consumption today to the worst consumption tomorrow, and that ratio is where the sandwich's slack of order human wealth enters — the same obstruction that closed the infinite-horizon routes in the companion write-up, met here for the fourth time and in the same form.

So the chain now has every link proved except its sharpness. Theorem 7.2 gives the criterion without risk, Theorem 7.4 signs the constraint's effect, Theorem 7.5 bounds it by a shift in wealth, and Theorem 7.6 converts a shift in wealth into consumption. What none of them supplies is a bound on consumption that is tight to a few per cent rather than to a factor of ten, and every route we have tried reduces to exactly that.

7.10 Carrying earnings across all remaining ages

The two halves of the sandwich are not equally to blame. The upper half, $c_k \leq \kappa_k(m + H_k)$, carries the expected present value of future earnings and is about twenty per cent above the truth at Aiyagari's calibration. The lower half, $\kappa_k m \leq c_k$, ignores future earnings entirely and is a factor of ten below it. The repair is the same device applied on the other side: carry earnings forward across all the remaining ages instead of one step at a time.

Theorem 7.7 (The lower half, with human wealth). *Let H_k be the present value of k periods of the minimum earnings, $H_0 = 0$ and $RH_{k+1} = y_{\min} + H_k$. Then $\lambda_k(m + H_k) \leq c_k(m)$ at every age and state, where $\lambda_0 = 1$ and $\lambda_{k+1} = p\lambda_k R/(p + \lambda_k R)$.*

Formal statement: *IncomeFluctuation.minHumanWealth_le_stageConsumption*.

The old bound is the case $H = 0$. The propensity λ_k is the finite-horizon propensity of Theorem 3.1 discounted once more by patience at each step, and the extra discount is not decoration: without it the bound is false at the borrowing constraint, where the household consumes exactly its cash on hand while $\kappa_k(m + H_k)$ exceeds it, because H_k discounts earnings

at R and $1/\kappa_k$ discounts at R/p . Discounting H_k at R/p instead repairs the corner and breaks the induction; the extra patience factor repairs both at once (`lowMPC_mul_minHumanWealth_1e`), and at $\beta = 0.96$, $r = 4\%$ and $\mu = 3$ it costs about three per cent over sixty periods.

At the state where the life cycle binds the household enters the next period with cash on hand about 2.0 and consumes about 0.88. The old bound gives 0.086; this one gives 0.33. Through Theorem 7.6 that improves the bound on the propensity to consume from about 0.94 to about 0.75, against a true propensity of 0.04.

So the multi-age argument buys a factor of four and not a factor of ten, and the reason is visible in the statement: it carries the *minimum* earnings, and a household that drew the minimum forever really would consume this little. What the programme needs is a lower bound carrying the *expected* human wealth, and that is not available from any chain of Euler inequalities, because the Euler inequality at a constrained state is exactly the statement that expected future earnings cannot be borrowed against. That is the one quantitative fact about a single household on which the whole programme turns, and identifying it precisely is what makes it attackable.

7.11 The constraint-incidence bound

Both halves can be carried at once. Accumulate human wealth by the *risk-adjusted* expectation — the power mean with exponent $-\mu$, which is what the Euler equation itself produces — and then *cap* it state by state at the level the borrowing constraint permits:

$$H_0 = 0, \quad H_{k+1}(z) = \min \left\{ \frac{1}{R} M_{-\mu}^z(y_w + H_k(w)), \left(\frac{1}{\kappa_{k+1}} - 1 \right) y_z \right\}$$

(`riskHumanWealth`), with κ the propensity of Theorem 3.1.

Theorem 7.8 (The constraint-incidence bound). $\kappa_k(m + H_k(z)) \leq c_k(a, z)$ at every age and state.

Formal statement: `IncomeFluctuation.riskHumanWealth_le_stageConsumption`.

The two ingredients meet exactly where they should. The power mean is what the Euler inequality delivers, and the step needs only that it is superadditive under a shift, which follows from its concavity and homogeneity (`powerMean_add_const_1e`). The cap is the constraint's incidence: at a corner the household consumes its cash on hand, which is at least y_z , and the cap is precisely the largest human wealth the bound may claim without exceeding that. No extra patience factor is needed, because the cap does the work the patience factor did in Theorem 7.7.

At Aiyagari's calibration, in the first period of a sixty-period life (`numerics/incidence.m`), the cap binds at the two lowest earnings states and nowhere else, which is where a household at zero assets is against the limit. The bound reaches

At zero wealth	lowest earnings	middle	highest
Bound as a fraction of the truth	1.000	0.890	0.877
Minimum-earnings bound, for comparison	1.010	0.409	0.295

and over the whole grid it never exceeds the truth, touching it exactly at the constrained corner. So the consumption bound is now within about twelve per cent where the life cycle binds, against a factor of three before. That is the sharpest bound in the development, and it was reached by bounding where the constraint binds rather than by another Euler chain.

It does not improve the propensity bound of Theorem 7.6, because that bound is driven by the worst earnings state next period, which is exactly where the cap binds. Closing the threshold argument would need the same sharpening on the worst state, and that turns out to be impossible within this class.

7.12 The worst state is pinned

Two facts settle it. The first is that the cap is not conservatism but necessity.

Theorem 7.9 (The cap is necessary). *Wherever a household at zero assets is against the borrowing limit, every affine lower bound $\kappa_k(m+H)$ valid at that state obeys $H \leq (1/\kappa_k - 1)y_z$.*
 Formal statement: *IncomeFluctuation.cap_necessary_of_corner.*

The proof is one line — at the corner consumption is exactly y_z — and the content is that Theorem 7.8 claims the largest human wealth any bound of its shape may claim. It cannot be sharpened within the affine class at all.

The second is that dropping the cap does not merely weaken the proof, it makes the bound false. Measuring effective human wealth as $c/\kappa - m$ at next-period assets of 1.67 (`numerics/worst.m`):

Effective human wealth at	truth	capped mean	uncapped mean	expectation
lowest earnings	10.96	5.99	13.11	16.57
highest earnings	27.92	23.90	26.38	32.24

At the lowest earnings state the uncapped power mean, 13.11, *exceeds* the truth, 10.96, so a bound built on it is invalid: the household will be against the limit later, and the uncapped recursion cannot see that. The truth sits between the two, and the distance is the price of capping at zero assets when the household in fact holds some. A bound that closed it would have to let human wealth depend on the assets held, which requires a lower bound on the household's own asset path; the one available, from the upper sandwich, is negative at the state that matters.

So the worst state is pinned from both sides: the cap cannot be raised, because the corner forbids it, and it cannot be dropped, because the truth is below what dropping it would claim. That closes this family of arguments, and with it, for now, the route from a consumption bound to the threshold. What the programme has instead is an exact statement of what a sharper argument must supply: a lower bound on consumption that is asset-dependent rather than affine, and the asset path to support it.

7.13 The bound with assets in it

The obstruction just described is no longer binding, because the path bound it asked for is available. Proposition 6.9 read at the household's own earnings state rather than at the worst

one is positive exactly where the cap bites hardest: a household in a high earnings state saves, so a period later it is not at the corner, and the cap that then applies to it is the one for a household with assets.

So assets go through the recursion twice over. Writing $\Phi_k(z, a)$ for that saving floor, clamped at zero,

$$G_0(z, a) = 0, \quad G_{k+1}(z, a) = \min\left(\frac{1}{R}M_{-\mu}^z(y_w + G_k(w, \Phi_{k+1}(z, a))), \left(\frac{1}{\kappa_{k+1}} - 1\right)(y_z + Ra)\right),$$

and the bound is $\kappa_k(m + G_k(z, a)) \leq c_k(a, z)$. The cap now grows with assets, and tomorrow's human wealth is read at the assets the household will hold rather than at zero.

Theorem 7.10 (The asset-dependent incidence bound). $\kappa_k(m + G_k(z, a)) \leq c_k(a, z)$, and $G_k(z, a) \geq H_k(z)$ for $a \geq 0$, so the bound is never worse than Theorem 7.8.

Formal statement: *IncomeFluctuation.assetHumanWealth_le_stageConsumption.*

, *IncomeFluctuation.riskHumanWealth_le_assetHumanWealth*

Measured against the truth over the whole grid at three ages (`numerics/assetdep.m`, `numerics/worstgrid.m`):

Worst ratio of bound to truth	$r = -7\%$	$r = 4\%$	$r = 10\%$
Theorem 7.8, cap at zero assets	0.237	0.606	0.760
Theorem 7.10, cap at the assets held	0.350	0.746	0.857

The new bound never exceeds the truth anywhere on the grid. The gain is largest exactly where the old one was weakest, at positive assets in low earnings states: at $r = 4\%$ with five units of assets and the lowest earnings, the ratio rises from 0.621 to 0.835. Where the old bound was already good, at zero assets, little changes, which is as it should be, since that is the state the cap was computed for.

What the refinement does not improve is the ceiling on aggregate capital, because the state that binds there is the highest earnings one, where the cap was never active: the implied ceiling at $r = -7\%$ is 22.15 against 22.17. That is still far below the 39.66 that feasibility alone gives, which is worth carrying into the existence argument of Section 6.

8 Theorem 1 at logarithmic utility

The slack condition of Section 6 is an Euler condition, and it is worth asking whether it settles the one case the stationary model does settle, $\mu = 1$. It does not, and the reason is instructive. Checking it on the whole grid at $\gamma = 1$ (`numerics/slacklog.m`), it holds at $r = 0$ and $r = 2\%$ and fails at $r = 4\%$ and above, by about twelve per cent in elasticity terms, at the youngest age with zero assets in the highest earnings state. Since a life-cycle equilibrium sits above the rate of time preference, 4.17% here, the Euler reduction says nothing where it is wanted, even at logarithmic utility.

8.1 A condition on the value gap

A different sufficient condition does hold. Write $V_k(\cdot; R)$ for the stage- k value.

Definition 8.1 (Value gap). The value gap has slope at least θ at stage k if, for $x \leq y$ in $[0, \bar{a}]$,

$$\theta(y - x) \leq [V_k(y; R_2) - V_k(y; R_1)] - [V_k(x; R_2) - V_k(x; R_1)].$$

Formal statement: `IncomeFluctuation.ValueGap`.

By the envelope theorem the slope of the gap is $R_2 u'(c_2) - R_1 u'(c_1)$, so the condition says the marginal value of assets rises with the rate. At $\gamma = 1$ that reads $c(R_2)/R_2 \leq c(R_1)/R_1$: consumption rises by proportionally less than the rate does.

Theorem 8.2 (Theorem 1 from the value gap). *If the value gap has a positive slope at stage k , then saving at stage k rises with the rate, at every state. Aggregate capital rises with it.*

Formal statement: `IncomeFluctuation.stagePolicy_mono_of_valueGap`.

, `IncomeFluctuation.olgCapital_le_of_valueGap`

Unlike the Euler route this needs no induction over ages. The condition on the continuation gives the conclusion for the household that faces it, directly, and the argument is Topkis with two continuations: the two households differ only in their cash on hand, which is ordered by the rate, and in their continuations, whose difference is increasing, so the higher-rate household cannot save less. Only concavity of period utility is used, as in Section 5.

8.2 Why it holds at $\mu = 1$, and only near there

At $\gamma = 1$ the patience factor is $\mathfrak{p} = \beta R$, so $q = \mathfrak{p}/R = \beta$ and the propensity $\kappa_k = 1/\sum_{i \leq k} \beta^i$ of Proposition 6.7 does not depend on the interest rate at all. Human wealth does, and falls as the rate rises. Where the borrowing constraint is slack, $c = \kappa_k(\bar{y} + Ra + H)$ and

$$R_2 c(R_1) - R_1 c(R_2) = \kappa_k[(R_2 - R_1)\bar{y} + R_2 H_1 - R_1 H_2] > 0,$$

which is the condition, with room. At a corner, $c = \bar{y} + Ra$ and the same inequality reduces to $R_1 \bar{y} \leq R_2 \bar{y}$.

Numerically (`numerics/condB.m`) the condition holds at $\gamma = 1$ at every age, state and rate tested, and fails from $\gamma = 2$, where between 19% and 51% of states violate it. It is therefore sufficient and not necessary: Theorem 1 itself still holds at $\gamma = 2$, and up to about 3.4. So this route covers $\mu = 1$ and closes nothing above it.

What is proved, then, is the reduction, and it is a reduction to a condition that is true at logarithmic utility rather than one that is false there. Discharging the condition itself from primitives is not available from the sandwich of Theorem 3.1: the sandwich's width is $\kappa_k H_k$, and the rate response of exactly that quantity is what drives the gap, so the bounds cannot see it. That is the same obstruction the consumption-bound family ran into above, and it is the piece a sharper argument would have to supply.

9 Mean field games, with age as time

The Lasry–Lions argument (Lasry and Lions, 2007) proves uniqueness in a mean field game from two optimality inequalities and a monotonicity condition on the cost. An earlier survey ruled it out for the stationary Bewley model, and for a stated reason: the argument cancels a term $\int (u_1 - u_2) d(m_1(0) - m_2(0))$ by assuming the two equilibria start from a common initial distribution, and a stationary discounted equilibrium has no common start — two candidate rates carry two different stationary distributions.

Overlapping generations do have one. Every cohort is born with nothing, whatever the rate, so the date-zero marginal is the same point mass crossed with the invariant earnings law at every candidate equilibrium. The stated obstruction is simply absent, and the monotonicity condition becomes a question to check rather than one to dismiss.

9.1 The game

Take age as time. A player is a cohort member, a path is $x = (a_0, \dots, a_J)$ with $a_0 = 0$, and the population is a distribution η over paths with date- t marginals m_t . The cost at date t is the negative of flow utility, with the budget substituted in,

$$F_t(x, m) = -\beta^t u\left(w(K(m))y_t + (1 + r(K(m)))a_t - a_{t+1}\right), \quad K(m) = \int a \, dm,$$

so the population enters through the single scalar K , and through it through both factor prices. This is the price-through-an-aggregate structure of Graber and Matter (2024), except that the aggregate enters the budget rather than being added to the cost.

9.2 What the argument actually uses

Adding the two equilibrium inequalities — each equilibrium does no worse against its own costs than the other distribution does — signs the sum of the date pairings and nothing finer:

$$\sum_t \underbrace{\int (F_t(\cdot, m_{1t}) - F_t(\cdot, m_{2t})) d(m_{1t} - m_{2t})}_{P_t} \leq 0.$$

Theorem 9.1 (The engine). *Two relaxed equilibria of the finite-horizon game with a common initial distribution have $\sum_t P_t \leq 0$. No monotonicity, no integrability and no continuity is used.*

Formal statement: `MeanFieldGame.summedPairing_nonpos_of_pathEquilibrium`.

The usual hypothesis, that every F_t is Lasry–Lions monotone, is one way to conclude from this and not the only one: it forces each summand to be zero. The weaker hypothesis that the argument can carry asks only that the total be nonnegative.

Theorem 9.2 (Uniqueness of an aggregate from the summed pairing). *If the cost separates a statistic A of the marginals — $\sum_t P_t > 0$ whenever A differs — then two relaxed equilibria from a common initial distribution agree on A . No single date is assumed monotone.*

Formal statement: *MeanFieldGame.aggregate_eq_of_pathEquilibrium*.

Date-by-date monotonicity is the special case in which every summand is already non-negative (*summedLasryLionsMonotone_of_forall*). The distinction is not cosmetic here, as the next two subsections show: in this economy the date-by-date hypothesis is false and the summed one is not.

9.3 The price channel is monotone; the curvature is not

Suppose for a moment that u were affine. Then $F_t(x, m_1) - F_t(x, m_2) = -\beta^t[(w_1 - w_2)y_t + (r_1 - r_2)a_t]$ and the pairing is

$$P_t = -\beta^t(w_1 - w_2) \int y_t d(m_{1t} - m_{2t}) - \beta^t(r_1 - r_2) \int a_t d(m_{1t} - m_{2t}).$$

The first term is exactly zero: labour is exogenous and the earnings marginal at each age is the same invariant law in both equilibria. The second is $-\beta^t(r_1 - r_2)(A_{1t} - A_{2t})$, and the factor-price frontier makes r strictly decreasing in K , so it is strictly positive at every date at which the economy with more capital holds more assets. The interest channel is exactly the Lasry–Lions congestion term: more aggregate capital raises the cost of holding capital. The price channel of a production economy does not run the wrong way. It is the only part of the problem that is unambiguously right.

What breaks is the weighting. With u strictly concave the pairing is that same price term weighted by marginal utility, and the weights do not cancel, because the two economies choose different savings and so sit at different marginal utilities. Over all pairs of measures the damage is total: taking concave and convex envelopes in the aggregate shows that unrestricted monotonicity of this cost forces u to be affine. The question worth asking is therefore the restricted one — whether monotonicity holds over the distributions the model can actually generate.

9.4 Date by date it fails, and it fails at birth

Take two economies differing only in β , each at its own equilibrium, so that each distribution is priced by the aggregate it actually carries. Since β does not appear in F_t , this is a legitimate test of the date- t condition. At the calibration of Section 6 ($\alpha = 9/25$, $\delta = 2/25$, $J = 60$, Aiyagari’s earnings chain), with $\beta = 0.960$ against $\beta = 0.950$:

γ	0.5	1	2	3	5
dates with $P_t < 0$ (of 60)	3	4	5	6	6
P_0	−0.00147	−0.00227	−0.00324	−0.00419	−0.00804
$\sum_t P_t$ (undiscounted)	+0.271	+0.175	+0.091	+0.061	+0.047

The failures are at the start of life, and they survive as the two equilibria close up: $P_0/(\Delta\beta)^2$ converges to −25.6 at $\gamma = 1$ and −44.4 at $\gamma = 3$ as $\Delta\beta$ falls from 0.02 to 0.00125,

with the count of negative dates unchanged throughout. So this is a property of the derivative, not an artefact of a wide gap.

Date zero is the sharp case, and it is worth spelling out, because it is exactly the common initial distribution that was supposed to rescue the argument. Every cohort is born with nothing, so $a_0 = 0$ for all of the mass in both economies, and the affine-utility pairing above is identically zero: with no assets to price, the aggregate coupling has no grip at all. But the wage still differs, and the utility gain from the higher wage is weighted by marginal utility, which is higher in the economy that saves more. At $\gamma = 1$ mean date-zero saving is 0.2937 against 0.2736, and P_0 is strictly negative. The common start does not make the date-zero pairing vanish; it removes the one term that could have made it positive.

9.5 Summed, it holds where the equilibrium lives

The summed condition is a different matter. Test it over the family the equilibria belong to: $m(\hat{K})$ generated by optimal life-cycle behaviour at the prices of $\hat{K}$, priced inside F by the aggregate $A(\hat{K})$ it actually carries. Every equilibrium is an $m(\hat{K})$ with $A(\hat{K}) = \hat{K}$, so monotonicity over pairs from this family is exactly what Theorem 9.2 needs.

γ	1	2	3	5
pairs with $\sum_t \beta^t P_t < 0$, $\hat{K}/K^* \in [0.5, 2]$	1/66	0/66	0/66	0/66
pairs failing inside $K \in [3.55, 7.40]$	0/36	0/36	0/36	0/28
least margin $\sum_t \beta^t P_t / (A_1 - A_2)^2$ there	+0.075	+0.097	+0.121	+0.338

The bracket $K \in [3.55, 7.40]$ is the one $r \in [2\%, 8\%]$ of Theorem 6.18, inside which the equilibrium is proved to lie. Within it the summed pairing is strictly positive for every pair and every risk aversion tested, with a margin bounded away from zero, and the margin grows with γ . The single failure is at $\gamma = 1$ between two candidates one of which sits outside the bracket entirely.

9.6 Towards a proof from primitives

The summed condition is an integral inequality over pairs of measures. It is not: it is a submodularity condition on a function of two real arguments, and the reduction is exact.

Let $U_i(K)$ be the lifetime utility of the *plan* carried by m_i — its asset path — when that plan is priced at aggregate K . Expanding the definitions and collecting terms,

$$\sum_t P_t = - \left[U_1(K_1) - U_2(K_1) - U_1(K_2) + U_2(K_2) \right],$$

a mixed difference. (Checked numerically: the two sides agree to 10^{-14} in every case reported above.) So summed Lasry–Lions monotonicity is exactly the statement that $U(K, \text{plan})$ has *decreasing differences* in the aggregate and the plan: the plan that goes with the larger aggregate loses its relative advantage as the aggregate rises.

Theorem 9.3 (The reduction). *When the population enters the cost only through a real statistic θ , the summed pairing is the mixed difference of the whole-life cost in (population, statistic). Strictly increasing differences of that cost make θ the same at every relaxed equilibrium with a common start.*

Formal statement: *MeanFieldGame.famSummedPairing_scalar.*

, MeanFieldGame.statistic_eq_of_increasingDifferences

This needed one repair. A stationary overlapping-generations aggregate is an average across *all* ages, so the date- t cost depends on the whole family of marginals and not on the date- t marginal alone. The finite-horizon machinery has been redone in that generality (*famPathCost*, *IsFamPathEquilibrium*, *famSummedPairing*, *famSummedPairing_nonpos_of_pathEquilibrium*); the engine survives unchanged.

What the differences are made of

Differentiate. For *any* constant-returns technology, Euler's theorem gives $\partial w / \partial K = -(K/L) \partial r / \partial K$, so with L normalised to one,

$$\frac{\partial U}{\partial K}(K, \text{plan}) = r'(K) \Phi(K, \text{plan}), \quad \Phi(K, \text{plan}) = \sum_t \beta^t \mathbb{E}[u'(c_t)(a_t - Ky_t)],$$

where c_t is the plan's consumption *at the pricing aggregate* K . The quantity $a_t - Ky_t$ is the household's **net capital position**: its assets less the capital that goes with its own labour at the economy's capital-labour ratio. Since $r' < 0$, decreasing differences say

$$\Phi(K, \text{plan}_1) \geq \Phi(K, \text{plan}_2) \quad \text{at every common } K, \quad (\text{W})$$

for plan_1 the plan carrying the larger aggregate: *the plan with more capital has the larger marginal-utility-weighted net capital position.*

A pointwise condition, and why it is no good

Submodularity of the flow $u(w(K)y + (1 + r(K))a - a')$ in (K, a) is, for constant relative risk aversion, exactly

$$\gamma(1 + r(K))(a_t - Ky_t) \leq c_t \quad \text{at every age and state,} \quad (\text{P})$$

which implies (W) and hence the summed condition. It is free wherever the household holds less than its share of the capital stock, and it is hopeless where it does not: at the calibration it is violated on 27% to 45% of the mass, at every risk aversion. Requiring each date to carry itself is precisely what the summed condition was introduced to avoid, and (P) is that mistake made again one level down.

The weighted condition holds where the pointwise one dies

(W) is a single scalar comparison, and it is weaker than the pointwise ordering of the two plans' asset paths. Measured at the calibration, on plans drawn from the candidate family at $\hat{K}/K^* = 0.90$ against 1.10 — so the two plans face different wages as well as different rates, which is the comparison (W) itself involves:

γ	1	2	3	5
$\Phi_1 - \Phi_2$ at the midpoint K	+39.36	+12.67	+7.34	+7.32
mass on which $a_t^1 \geq a_t^2$ FAILS	0.0000	0.0025	0.0409	0.2006
the same, at a common wage	0	0	0	0.0091

The second row is the one to read against Light's Theorem 1, which is a statement about the household at a *fixed* wage: in wage units the policy ordering is exact — violated on zero mass, not on small mass — for $\gamma \leq 3$, and fails on under one percent of the mass at $\gamma = 5$ (on 6.7% at $\gamma = 8$). The first row, which lets the wage move along the frontier as well, overstates the failure of Theorem 1 by more than an order of magnitude; it is the right measure for (W) and the wrong one for Theorem 1. So (W) is weaker than the pointwise ordering and does survive where that ordering fails, but the margin between them at $\gamma = 5$ is one percent of the mass, not a fifth.

More than those pairs is true. Sweeping the planning aggregate over $\hat{K}/K^* \in [0.5, 2]$ at fourteen points, $\Phi(K, p(\hat{K}))$ is strictly *decreasing* in $\hat{K}$ at every one of the thirteen steps, for $\gamma = 1, 3, 5$ and for each of $K = 0.90K^*, K^*, 1.10K^*$. So (W) holds not merely at the pairs one happens to test but as a monotone comparative static in the plan. Over all 91 pairs at five pricing points each, $\gamma = 3$ and $\gamma = 5$ give no failures at all; the 42 failures out of 455 at $\gamma = 1$ all involve the extreme candidate $\hat{K} = 0.5K^*$, whose carried aggregate is 48.6, ten times the capital stock and far outside the bracket of Theorem 6.18.

9.7 What Φ is, and why it will not yield locally

Two things can be proved. Neither closes (W), and the second says why.

The duration identity

Along the diagonal, where a plan is priced by the aggregate it was chosen for, Φ has a closed form. Let $m_t = \beta^t u'(c_t)$. Interiority makes m a state-price density, $\mathbb{E}_t[m_{t+1}] = m_t/R$, and writing $P[x] = \sum_t \mathbb{E}[m_t x_t]$ and $P_1[x] = \sum_t t \mathbb{E}[m_t x_t]$, two identities follow from the budget recursion $a_{t+1} = Ra_t + wy_t - c_t$ with $a_0 = a_{J+1} = 0$: the lifetime budget in state-price units, $P[wy - c] = 0$, and $P[a] = -P_1[wy - c]/R$.

Proposition 9.4 (Φ as a duration wedge). *For the plan optimal at K , with $Y = P[y]$ and $T_c = P_1[c]/P[c]$, $T_y = P_1[y]/P[y]$ the marginal-utility-weighted durations of consumption and of earnings,*

$$\Phi(K, p(K)) = Y \left[\frac{w}{R} (T_c - T_y) - K \right].$$

Assets are the wedge between the duration of spending and the duration of earning; the household's marginal valuation of a higher interest rate is negative exactly when that wedge, scaled by the wage, falls short of the economy's capital-labour ratio. At the calibration it does, at every risk aversion: $T_c - T_y$ runs from 1.7 to 3.8 years while $K \approx 5$, so $\Phi < 0$ — along the factor-price frontier a higher interest rate *hurts* the household, because $dw/dr = -K$ and the marginal-utility weighting favours the young, who have no assets to gain on. Checked numerically: $P[wy - c]/P[c] = 3 \times 10^{-5}$, $P[a] + P_1[wy - c]/R = 1.7 \times 10^{-2}$ against $P[a] = 57.9$, and the two sides of Proposition 9.4 agree to one part in 10^4 ; the residual is the borrowing constraint, which binds on 3–11% of the mass and breaks the Euler equality there. At $\gamma = 1$ the identity $m_t c_t = \beta^t$ makes $T_c = \sum_t t \beta^t / \sum_t \beta^t$ exactly, independent of prices and of the earnings process.

This is the diagonal, and (W) is about two plans priced at a *common* aggregate that is neither's own. The identity does not reach it.

The obstruction

The reason (W) resists is not that the argument is delicate. It is that the inequality is the whole of uniqueness, and the following says so.

Proposition 9.5 (The pairing is minus two revealed-preference losses). *For any cost, the summed pairing satisfies the identity $\sum_t P_t = -[(C_1[m_2] - C_1[m_1]) + (C_2[m_1] - C_2[m_2])]$, where $C_i[m]$ is the whole-life cost borne by population m at the prices set by population i . At two relaxed equilibria with a common start both brackets are nonnegative.*

Formal statement: `MeanFieldGame.famSummedPairing_eq_neg_losses`.

`, MeanFieldGame.famCost_le_of_pathEquilibrium, MeanFieldGame.famSummedPairing_nonpos_of_losses`

At two *distinct* equilibria both brackets are strictly positive, so $\sum_t P_t < 0$ strictly. A hypothesis forcing $\sum_t P_t > 0$ is therefore not a statement about nearby populations that could be established by a local expansion and then extended: it is, on the nose, the assertion that no second equilibrium exists. That kills the route we would otherwise have taken — expand the optimality loss to second order about the diagonal, where it is a positive quadratic, and read off a radius within which the sign is guaranteed. There is no such radius, because the sign that a local expansion delivers at a hypothetical pair of equilibria is the wrong one. (W) is strictly stronger than uniqueness, as Light's Theorem 1 also is, and any proof of it has to be global.

Optimality gives the differences falling for free over the interval between the two *planning* aggregates — add the two optimality inequalities — while what is needed is the same fall between the two *carried* aggregates; the intervals coincide only at equilibrium, and there Proposition 9.5 turns the free inequality against us.

9.8 The global attempt, and where it ends

Since no local argument can work, the proof has to be global. Pushing it there settles the route, though not in the direction we wanted.

Re-pricing a plan, exactly

One piece of structure is worth recording because it is exact and it is the natural starting point for any global argument. Parametrise by the rate rather than the aggregate. The factor-price frontier is $dw/dr = -K(r)$, so for a plan with asset path a chosen at rate $\hat{r}$ and re-priced at rate r ,

$$w(r) - w(\hat{r}) = \int_{\hat{r}}^r w'(s) ds = -(r - \hat{r}) \bar{K}, \quad \bar{K} := \frac{1}{r - \hat{r}} \int_{\hat{r}}^r K(s) ds,$$

and therefore, at every date and in every state,

$$c_t(r) = \hat{c}_t + (r - \hat{r})(a_t - \bar{K}y_t). \quad (\text{R})$$

Re-pricing shifts consumption by the rate gap times the net capital position *at the average capital intensity over the interval*. It is an identity, not an approximation. One consequence is immediate: differentiating Φ in the pricing gap at a fixed plan gives $\sum_t \beta^t \mathbb{E}[u''(c_t)(a_t - \bar{K}y_t)(a_t - Ky_t)]$, which is negative wherever the two net positions agree in sign — everywhere outside a band of asset-to-earnings ratios of width $|K - \bar{K}| = O(r - \hat{r})$. So Φ falls in the pricing gap at a fixed plan. What (W) needs is the comparison across *different* plans, and (R) does not supply it.

The weakest form of the hypothesis

Before grinding further it is worth asking what the least one could get away with is. The uniqueness chain runs: two stationary equilibria (K_1, r_1) , (K_2, r_2) with $K_1 > K_2$; the firm's demand is strictly decreasing, so $r_1 < r_2$; the hypothesis gives $\Delta^2 \leq 0$; revealed preference gives $\Delta^2 > 0$ by Proposition 9.5; contradiction. Nothing in that chain asks capital supply to be monotone in the rate. All it asks is

$$A(\hat{r}_1) > A(\hat{r}_2) \implies \int [\Phi(\cdot, p(\hat{r}_1)) - \Phi(\cdot, p(\hat{r}_2))] \geq 0, \quad (\text{C})$$

that the carried aggregate A and the weighted net position Φ be *comonotone* along the family of optimal plans. Neither is required to be monotone in its own right. That is genuinely weaker in kind than the single-crossing route, which needs A itself increasing — the step this programme cannot take above $\mu \approx 3.4$.

The collapse

It is weaker in kind, and it is not weaker here.

Proposition 9.6 (Comonotonicity collapses to monotone supply). *If A and Φ are comonotone along the family and Φ is strictly monotone along it, then A is monotone along it.*

Formal statement: `MeanFieldGame.monotone_of_comonotone_of_strictMono`.

And Φ is strictly monotone along the family: that is the measurement of Section 9, at every one of thirteen steps across $\hat{K}/K^* \in [0.5, 2]$, at every pricing point and every risk aversion. So (C) delivers monotone capital supply — and monotone capital supply already gives uniqueness by single crossing, with no mean field game anywhere in the argument. The Lasry–Lions route, taken to its end, is not logically weaker than the route it was meant to bypass.

This is a definite closure rather than an open problem, and it is worth being exact about what it does and does not say. It does not say (W) is false: every measurement says it is true. It says that a proof of (W) would be a proof of monotone capital supply, so there is no shortcut hiding in the mean field formulation. The one escape is an economy in which Φ is *not* monotone along the family, where (C) would be strictly weaker than monotone supply; we have not found one, and at this calibration none exists.

9.9 What this leaves

Three things, in order of how much they are worth.

The negative result is definite and it is not the expected one. Date-by-date Lasry–Lions monotonicity fails in this economy, for every risk aversion including logarithmic, and the reason is not that the price channel runs backwards — it runs the right way, exactly — but that strict concavity reweights it, most severely at the ages where there are no assets to price. This is not a $\mu = 1$ frontier of the kind the rest of the programme keeps running into; it fails at $\mu = 1$ too.

The structural gain is real. Theorem 9.1 isolates what the argument does without any hypothesis, and Theorem 9.2 turns the summed condition directly into uniqueness of the aggregate, which for this economy is uniqueness outright: one K means one r , one household problem, one answer. The route is no longer dismissed for the wrong reason.

What is missing is a proof of (W) from primitives, and Proposition 9.6 says why it will not come cheaply: a proof of it would be a proof of monotone capital supply. The condition is strictly weaker than the *pointwise* policy ordering — it survives at $\mu = 5$ where that fails, and its measured margin improves as risk aversion rises — but it is not weaker than the aggregate monotonicity that the single-crossing route needs. As a route to uniqueness this one is closed. What survives it is the reduction itself: the summed pairing is a mixed difference (Theorem 9.3), Φ is a duration wedge (Proposition 9.4), and re-pricing is exact (R) — all of which are statements about the life-cycle household rather than about mean field games, and none of which depended on the route working.

10 The increment of capital supply

Section 9 ends by closing a route and leaving the original target: uniqueness through single crossing, which the programme has pursued through Light's Theorem 1 and cannot get above $\mu \approx 3.4$. This section changes what has to be proved. Theorem 1 need not hold. It need only fail on little mass.

10.1 What single crossing actually asks

Work in wage units, where homotheticity makes capital supply per unit of wage a function of the rate alone and clearing reads $S(r) = D(r)$ with $D(r) = \alpha/((1 - \alpha)(r + \delta))$. Excess supply is increasing — which is all single crossing needs — as soon as

$$S(r_2) - S(r_1) > D(r_2) - D(r_1), \quad r_1 < r_2, \quad (\text{I})$$

that is, as soon as supply does not fall *faster* than demand. Demand is steep at the calibration, and that is the whole opportunity: over the proved bracket $r \in [2\%, 8\%]$ supply's worst secant slope is +33 at $\gamma = 5$ against -54 for demand, and it never bends backwards at any γ up to 12. A proof of (I) may be lossy by more than the true increment and still close.

The two-sided bracket on the *level* of capital that the existence theorem supplies cannot be used for this. As $r_2 - r_1 \rightarrow 0$ the bracket-only condition $F(r_2) - C(r_1) > D(r_2) - D(r_1)$ degenerates into $F \geq C$, a bracket of zero width; measured, its margin runs from -352 to -1737 . The object to bound is the increment, not the level twice.

10.2 Coupling the two economies

Run both economies on a common earnings path from the common start $a_0 = 0$, and let $x_j = a_j(r_2) - a_j(r_1)$. Split the step at the *lower*-rate state:

$$x_{j+1} = \underbrace{[g_j^2(a_j^2) - g_j^2(a_j^1)]}_{\text{same policy, different assets}} + \underbrace{[g_j^2(a_j^1) - g_j^1(a_j^1)]}_{\text{same state, different rate}}.$$

The second bracket is the failure of Theorem 1 at that state, which (`stagePolicy_le_iff_consumption_le`) is exactly $c^2 - c^1 \leq (r_2 - r_1)a$ read with a slack; call its shortfall e_j . The first is at least $-L_j(x_j)^-$ when saving is nondecreasing in assets and Lipschitz with constant L_j . So

$$x_{j+1} \geq -L_j(x_j)^- - e_j. \quad (*)$$

Nothing in $(*)$ asks $x_j \geq 0$: the pointwise ordering of policies is not assumed, and its failure is priced by e_j rather than excluded.

Theorem 10.1 (The gap envelope). *If $x_0 \geq 0$, $L_j \geq 0$, $e_j \geq 0$ and $(*)$ holds at every age, then $x_j \geq -f_j$ where $f_0 = 0$ and $f_{j+1} = L_j f_j + e_j$; hence $\sum_j x_j \geq -\sum_j f_j$. With demand falling by more than the average of f , (I) follows.*

Formal statement: *OLG.gapAccum*.

, *OLG.negPart_le_gapAccum*, *OLG.neg_gapAccum_le*, *OLG.sum_neg_gapAccum_le*, *OLG.lt_of_increment_bound*

10.3 Why the bound is not worst-case

The negative part is the point of the formulation. From (*),

$$(x_{j+1})^- \leq L_j (x_j)^- + e_j,$$

whose right-hand side is *linear* in the two quantities. Expectations therefore pass straight through it, and averaging over the cohort replaces e_j by the **mass-weighted** failure of Theorem 1, $\bar{e}_j = \mathbb{E}_1[(g_j^1 - g_j^2)^+]$, in place of its worst case over states. Since the violation is confined to a sliver of the distribution, that is the difference between a bound that fails and one that passes by two orders of magnitude.

Measured at the calibration, with $L_j = (1 - \kappa_k)(1 + r_2)$ from the sandwich and the accumulation $\Gamma = \frac{1}{j} \sum_j f_j / \varepsilon$ around 21 to 24:

γ	$r \in [4\%, 6\%]$, budget 0.670		$r \in [2\%, 8\%]$, budget 2.109	
	5	8	5	8
worst violation ε	0.0460	0.0701	0.1313	0.2053
mass carrying it	0.0091	0.0666	0.0084	0.0636
bound, uniform ε	$0.976 \times$	$1.464 \times$	$3.125 \times$	$4.707 \times$
bound, age-dependent ε_j	$0.760 \times$	$1.218 \times$	$2.498 \times$	$3.953 \times$
bound, mass-weighted $\bar{e}_j$	0.0046 ✓	0.0310 ✓	0.0138 ✓	0.0958 ✓
true $S(r_2) - S(r_1)$	+0.692	+0.457	+2.060	+1.368

At $\gamma \leq 3$ every column is zero: in wage units Theorem 1 holds exactly, so $\varepsilon = 0$ and the bound is trivial. The interest is at $\gamma = 5$ and 8, where Theorem 1 fails and the mass-weighted bound passes with factors of 145 and 22 to spare — and passes on the wide bracket too, where the rates differ by six percentage points. Charging the worst violation at every age fails; charging it where it occurs does not. Age-dependence alone does not rescue the worst-case version, because the violation is spread over 53 of 60 ages while being tiny at nearly all of them.

10.4 What remains

The engine is proved, and so is the Lipschitz constant it runs on.

The propensity floor

Theorem 10.2 (The propensity to consume is at least κ_k). *At every stage, every income state and every pair of asset levels in the region,*

$$\kappa_k (1 + r)(a_2 - a_1) \leq c_k(z, a_2) - c_k(z, a_1), \quad a_1 \leq a_2,$$

so saving is Lipschitz in assets with constant $(1 - \kappa_k)(1 + r)$, which is the L_j the increment argument uses.

Formal statement: *IncomeFluctuation.stageMPC_mul_le_stageConsumption_sub*.

The terminal condition is an identity rather than an assumption: in the last period consumption is cash on hand, so the propensity is exactly $1 = \kappa_0$. The step is the two-point Euler comparison. With b_i the saving at a_i and $t = \kappa_k R(b_2 - b_1) \geq 0$, the induction hypothesis raises next period's consumption by at least t in *every* state, so the Euler equation above the floor at a_2 gives

$$c_2^{-\gamma} \leq \beta R \sum_{z'} \pi(C'(z') + t)^{-\gamma} = \beta R M(C' + t)^{-\gamma} \leq \beta R (M(C') + t)^{-\gamma},$$

with M the certainty equivalent, the power mean of exponent $-\gamma$. The Euler equation under the cap at a_1 gives $M(C') \geq \mathfrak{p}c_1$, and $\beta R = \mathfrak{p}^\gamma$ turns the chain into $c_2 \geq c_1 + t/\mathfrak{p}$, which is the claim at stage $k + 1$ once denominators are cleared.

The last inequality in the display is **superadditivity of the certainty equivalent**, $M(x + t) \geq M(x) + t$, which is concavity and homogeneity of degree one at the midpoint (*powerMean_add_const_le*, already proved for the constraint-incidence bound of Section 7). It is worth pausing on: this is the direction a negative power mean obliges, and it is the reverse of what several arguments in this programme wanted and could not have. The recurring obstruction, read the other way round, is what makes the floor go through.

Measured, the floor holds at every stage and every reached state with equality only at the last stage, and the constant is nearly tight: the true accumulation is 20.7 against the sandwich's 21.2.

The Lipschitz bound it gives

Theorem 10.3 (Accumulated assets are Lipschitz in assets). *For $a_1 \leq a_2$ in the region,*

$$A_k(a_2, z) - A_k(a_1, z) \leq G_k(a_2 - a_1),$$

*with G_k the slope *cohortSlope* of the affine ceiling: $G_0 = 1$, $G_{k+1} = 1 + G_k(1 - \kappa_{k+1})(1 + r)$.*

Formal statement: *IncomeFluctuation.cohortAssets_sub_le_cohortSlope*.

The induction is the one already written for the affine floor, with one substitution: the propensity floor turns a step of assets into a step of saving no larger than $(1 - \kappa_{k+1})(1 + r)$ times it, and the expectation over next period's earnings passes the constant through unchanged. The same G_k that bounds the *level* of accumulated assets in Section 6 therefore bounds their *differences*, which is what the increment argument accumulates — and it is the L_j whose accumulation was measured at 21 above.

10.5 The mass-weighted failure

What is left is $\bar{e}_j$, an expectation against a distribution nobody controls. It looks like the hard part, and it would be if the failure of Theorem 1 were spread over assets. It is not.

Where the failure lives

Measured at $\gamma = 5$ over $r \in [4\%, 6\%]$, the failure reaches down to $a = 0$ — so no threshold in assets confines it, and a Markov bound on a thin upper tail is worthless: the sup of the failure below any threshold is the global sup, and the resulting bound is the uniform one that already failed. But sorted by earnings state the picture is different. With Aiyagari’s seven-state chain the sup of the failure over assets, state by state, is

$$(0, 0, 0, 0, 0, 0.0143, 0.0460),$$

zero in the bottom five states and confined to the top two, whose invariant mass is $\nu_6 + \nu_7 = 0.095$. At $\gamma = 8$ it reaches the fourth state and the profile is $(0, 0, 0, 0.0014, 0.0118, 0.0333, 0.0701)$. The failure of Theorem 1 is an *earnings-state* phenomenon: it is the high-earners, who are few.

Why that coordinate is free

Earnings are exogenous and the cohort is born with its state drawn from the invariant law, so the earnings marginal is ν at *every* age, exactly. Nothing has to be estimated in that coordinate.

Theorem 10.4 (The earnings marginal of a cohort is ν). *If ν is invariant for π , the ν -weighted expectation of a function of the earnings state is unchanged by any number of steps of the cohort recursion. Hence $\bar{e}_j \leq \sum_z \nu_z V_j(z)$ with $V_j(z)$ the supremum over assets of the failure at earnings state z : exact in earnings, worst-case only in assets.*

Formal statement: *IncomeFluctuation.incomeStep.*

, IncomeFluctuation.stageStep_of_incomeFn, IncomeFluctuation.sum_mul_incomeStep, IncomeFluctuation.sum_mul_incomeStep_iterate, IncomeFluctuation.sum_mul_le_of_le

Being exact in the coordinate where the failure lives is what makes the bound work:

γ	$r \in [4\%, 6\%]$, budget 0.670		$r \in [2\%, 8\%]$, budget 2.109	
	5	8	5	8
worst-case in both coordinates	0.976 ×	1.464 ×	3.125 ×	4.707 ×
Markov in assets, best threshold	0.976 ×	1.464 ×	3.125 ×	4.707 ×
ν -exact, worst-case in assets	0.032 ✓	0.130 ✓	0.103 ✓	0.422 ✓
fully mass-weighted (the truth)	0.0046	0.0310	0.0138	0.0958

The ν -exact row passes with factors of 21, 5.1, 20 and 5.0, and it is within a factor of seven of the fully mass-weighted truth — the whole of the remaining looseness being the worst case over assets inside each earnings state.

10.6 What $V_j(z)$ would take

One quantity remains: $V_j(z)$, the largest failure of Theorem 1 over assets at a single earnings state. There is no distribution in it, and the accumulation leaves an allowance of $0.670/21 \approx 0.032$ per age for $\sum_z \nu_z V_j(z)$. Two routes suggest themselves. Both fail, and the reasons are worth recording because they are quantitative.

The quantitative step amplifies

The Theorem-1 step can be made quantitative without touching a power mean. Suppose the failure $D = c_2 - c_1 - \Delta r a$ is positive at (a, z) . Then $b_1 > b_2 \geq 0$, so the low-rate Euler equation holds above the floor and the high-rate one under the cap; the stage- k failure bound and monotonicity of consumption in assets give $C_2(z', b_2) \leq C_1(z', b_1) + \Delta r b_1 + \varepsilon_k(z')$, and $(C + s)^{-\gamma} = C^{-\gamma}(1 + s/C)^{-\gamma} \geq C^{-\gamma}(1 + \sigma)^{-\gamma}$ with $\sigma = \max_{z'} s(z')/C_1(z', b_1)$. Chaining the two Euler relations,

$$D \leq c_1 \left[(R_1/R_2)^{1/\gamma} (1 + \sigma) - 1 \right] - \Delta r a. \quad (\text{Q})$$

This is a genuine quantitative step: it recovers the no-failure criterion when the bracket is nonpositive, and bounds the failure otherwise. It also amplifies. The derivative of (Q) in ε_k is $c_1(R_1/R_2)^{1/\gamma} / \min_{z'} C_1(z', b_1)$, measured at 4.7 at $\gamma = 5$ and 4.2 at $\gamma = 8$. Over sixty stages the recursion started from the exact terminal condition $\varepsilon_0 = 0$ reaches $V \approx 2.8 \times 10^3$ at $\gamma = 5$ against a true supremum of 0.046 — and with the sandwich floor in place of the true C_1 , 10^{60} . This is the amplification the rate-response route already ran into, with a per-step multiplier bounded below by $1/\text{p} > 1$; the life cycle does not escape it, it merely makes it finite.

The direct bound is an order of magnitude too loose

The alternative is to bound the failure at each stage separately, with no induction, from the sandwich: $c_2 \leq \kappa_k^2(m_2 + H_k^2)$ and $c_1 \geq \kappa_k^1 m_1$. That gives 1.835 against an allowance of 0.032, too loose by a factor of 58, and the source of the looseness is visible: the sandwich slack is of order $\kappa_k H_k$, measured at 0.96, an order-one quantity where the target is of order 0.03.

What that says

The remaining gap is not a matter of assembling existing bounds. It needs an estimate of the rate-response of consumption whose error is small *relative to the response itself* — an absolute error below 0.03 on a quantity of size about one — and no sandwich in this programme is of that quality. That is a sharper statement of the requirement than the programme has had before, and it is attached to a single distribution-free object rather than to Theorem 1 in full.

What is proved stands on its own: single crossing needs only the increment (I); the increment is controlled by the envelope of Theorem 10.1; its Lipschitz constant is Theorem 10.2 and Theorem 10.3; and the mass weighting is exact in the coordinate where the failure lives, by Theorem 10.4. Numerics: `incbound.m`, `ebar.m`, `vjbound.m`, `vjdiag.m`, `mpclow.m`, `slopecheck.m`.

11 The programme

The stages, in the order we intend to take them. The first is done (Section 3).

Equilibrium. A stationary overlapping-generations equilibrium: the distribution μ_1 of the youngest cohort is given (zero wealth, earnings drawn from the invariant law), μ_{j+1} is the pushforward of μ_j under the age- j policy and the earnings chain, capital supply is $K(r) = \sum_j \omega_j \int a d\mu_j$ with cohort weights ω_j , labour is exogenous, and the firm is Aiyagari's. Targets: continuity of K in the rate, a finite composition of continuous pushforwards; existence by the intermediate value theorem against the firm's demand, with the floor at the top of the interval from explicit accumulation over the life rather than from an ergodic identity; Light's Theorem 2 by induction on age, an ordering of policies at every age giving first-order stochastic dominance of every μ_j and so of K ; and single crossing against a decreasing demand.

Theorem 1 by backward induction. The research core: saving rises with the rate at every age, for $\mu \in \{3, 5\}$. The reduction of the companion write-up applies stage by stage, Theorem 1 at stage $k + 1$ following from the slack elasticity condition on the stage- k consumption functions at the pivot; the transfer lemma propagates an affine bound on the response of consumption to the rate backward from the terminal condition $c_0 = m$, where the response is exactly $\Delta r \cdot a$; and the slack condition at each stage is a finite check on the constants so obtained. The obstruction of the infinite horizon, a step multiplier of at least $1/p$, does not bite in J steps. The first session is numerical: a life-cycle economy with $J = 60$, Aiyagari's earnings chain and $\mu \in \{3, 5\}$, checking the slack condition at every age and wealth and running the recursion for the constants. If it holds the argument is formalised; if it fails only at moderate ages and wealth we look for the age-dependent refinement; if it fails badly, log utility is the fallback, where Theorem 1 at each stage follows from the rescaling proof of the companion, which is already an induction on Bellman iterates.

Diamond (1965) and Samuelson. Section 2 does the first item, and for every risk aversion rather than only for logarithmic utility. What remains of the two-period economy: monotone global convergence, existence under Inada conditions and a multiplicity witness in general, dynamic inefficiency of a steady state with $r < n$ by an explicit Pareto-improving transfer, the golden rule, and the autarkic and monetary steady states of the pure-exchange economy. These are standalone results and will be the classical sections of this write-up.

Uniqueness: bound the increment, not the level. Section 10 does this, and records what is left of it: the propensity-to-consume floor, the Lipschitz bound on `cohortAssets`, and an estimate of the mass-weighted failure of Theorem 1.

Uniqueness: what stochastic dominance does not buy. Light and Weintraub's Theorem 2 reaches uniqueness from ordered stationary distributions, and its hypothesis

delivers the ordering through the pointwise policy ordering. Since the policy ordering fails above $\gamma \approx 3.4$, the obvious weakening is to ask only for the ordering it was used to produce: first-order stochastic dominance of the cohort distributions. Measured, it buys nothing. In wage units, at $r = 4\%$ against 6% , dominance holds to machine precision at $\gamma \leq 3$ — where the policy ordering already holds exactly — and fails at precisely the risk aversions where the policy ordering fails: at $\gamma = 5$ eleven of sixty ages have a crossing of the asset CDF, worst gap 5×10^{-2} , and conditional on the income state the worst gap is 0.57 ; at $\gamma = 8$, nineteen ages and 0.22 . The distributions genuinely cross, and they stop being ordered exactly when the policies do. Aggregate supply stays monotone throughout ($4.54 \rightarrow 5.24$ at $\gamma = 5$, $5.53 \rightarrow 5.98$ at $\gamma = 8$), which is the point: the aggregate is the weaker statement, and there is no distributional route to it. Numerics: `fosdcheck.m`.

Extensions. An age profile of earnings, retirement with a pension and a government budget, survival risk with accidental bequests as a second fixed point, population growth, and Lasry–Lions with age as time, whose finite-horizon core is already formalised and whose monotonicity is again the elasticity condition.

References

- AIYAGARI, S. R. (1994): “Uninsured Idiosyncratic Risk and Aggregate Saving,” *Quarterly Journal of Economics*, 109, 659–684.
- DIAMOND, P. A. (1965): “National Debt in a Neoclassical Growth Model,” *American Economic Review*, 55, 1126–1150.
- GRABER, P. J. AND E. MATTER (2024): “New Uniqueness Results for a Mean Field Game of Controls,” ArXiv:2410.13950.
- HUGGETT, M. (1996): “Wealth Distribution in Life-Cycle Economies,” *Journal of Monetary Economics*, 38, 469–494.
- LASRY, J.-M. AND P.-L. LIONS (2007): “Mean Field Games,” *Japanese Journal of Mathematics*, 2, 229–260.

A Index of formal results

All in the namespace `LeanEconomics.IncomeFluctuation`, directory `LeanEconomics/OLG`. Every result depends only on the three standard axioms (`propext`, `Classical.choice`, `Quot.sound`) and none uses `sorry`.

Result	Lean name	Module
Stage value, policy, consumption	<code>stageValue</code> , <code>stageConsumption</code>	<code>stagePolicy</code> , <code>OLG/LifeCycle</code>

Result	Lean name	Module
Last period: no saving	stagePolicy_zero	OLG/LifeCycle
Last period: consume cash on hand	stageConsumption_zero	OLG/LifeCycle
Positive consumption at every stage	stageConsumption_pos	OLG/LifeCycle
Saving and consumption rise with wealth	stagePolicy_ mono, stageConsumption_mono	OLG/LifeCycle
Carroll–Kimball at every stage	concaveOn_stageConsumption	OLG/LifeCycle
Euler under the cap, by stage	stage_euler_le	OLG/LifeCycle
Euler above the floor, by stage	stage_euler_ge	OLG/LifeCycle
Finite-horizon propensities, human wealth	stageMPC, stageHumanWealth	OLG/LifeCycle
Propensities in $(0, 1]$, above κ	stageMPC_pos, stageMPC_le_one, minMPC_le_stageMPC	OLG/LifeCycle
Lower Euler step	crRa_stageMPC_step	OLG/LifeCycle
Upper Euler step	crRa_stageUpper_step	OLG/LifeCycle
Theorem 3.1, lower	stageMPC_mul_le_stageConsumption	OLG/LifeCycle
Theorem 3.1, upper	stageConsumption_ le_ stageHumanWealth	OLG/LifeCycle
One-step saving bound	stagePolicy_le	OLG/LifeCycle
Retired household (lump-sum pension)	withPension	OLG/Retirement
Cash on hand rises with assets	withPension_resources_strictMono	OLG/Retirement
Human wealth is state-independent	stageHumanWealth_ withPension_ const	OLG/Retirement
Perpetuity bound $H_k \leq p/r$	stageHumanWealth_withPension_le	OLG/Retirement
Life-cycle value with retirement	lifeValue	OLG/Retirement
Saving, consumption rise with wealth	lifePolicy_mono, lifeConsumption_mono	OLG/Retirement
Proposition 2.1	crRa_le_tangent	OLG/Diamond
Closed-form saving	saveShare, saveShare_euler	OLG/Diamond
It is the optimum	saveShare_isOptimal	OLG/Diamond
Proposition 2.2	saveShare_strictMonoOn, saveShare_strictAntiOn	OLG/Diamond
Proposition 2.3	div_one_add_rpow_strictMono	OLG/Diamond
Theorem 2.4	exists_unique_steadyState	OLG/Diamond
Theorem 4.1	stageConsumption_eq_of_interior	OLG/ Deterministic
Rate-free propensity at log	stageMPC_one_eq_of_rate	OLG/ Deterministic
Human wealth falls with the rate	stageHumanWealth_antitone_rate	OLG/ Deterministic
Theorem 4.2	stagePolicy_ mono_ det_ log, olgCapital_mono_det_log	OLG/ Deterministic
Chord of a concave power	CertaintyEquivalent.chord_le_rpow	OLG/Certainty Equivalent
Theorem 7.3	CertaintyEquivalent.consumption_antitone_of_chord	OLG/Certainty Equivalent
Pension schedules	lumpSum, tapered, cliff	OLG/MeansTest
Cash on hand	cash	OLG/MeansTest
Theorem 5.1	isOptimalSaving_mono	OLG/MeansTest

Result	Lean name	Module
Monotone in assets when cash is	<code>isOptimalSaving_mono_assets</code>	OLG/MeansTest
Gentle taper keeps cash rising	<code>cash_strictMono_tapered</code>	OLG/MeansTest
A cliff reverses cash on hand	<code>cash_lt_cash_of_cliff</code>	OLG/MeansTest
Theorem 5.2	<code>exists_consumption_decrease</code>	OLG/MeansTest
Theorem 5.3, cliff	<code>exists_saving_decrease_in_assets</code>	OLG/MeansTest
Theorem 5.4	<code>policyOf_mono_of_lt</code>	OLG/MeansTest
One age's expectation operator	<code>stageStep</code>	OLG/ Equilibrium
Monotone class preserved	<code>monoRegion_stageStep</code>	OLG/ Equilibrium
Ordered policies, ordered functions	<code>stageStep_le_stageStep</code>	OLG/ Equilibrium
Life-cycle assets, aggregate capital	<code>cohortAssets, olgCapital</code>	OLG/ Equilibrium
Theorem 6.1	<code>cohortAssets_le, olgCapital_le_of_</code>	OLG/ Equilibrium
Proposition 6.9	<code>stagePolicy_le</code>	OLG/ Equilibrium
	<code>le_stagePolicy</code>	OLG/ Equilibrium
Theorem 6.2	<code>olgEquilibriumRate_unique</code>	OLG/ Equilibrium
Argmax at a concave continuation	<code>argmax_eq_singleton_of_</code>	OLG/Continuity
	<code>concaveSlices</code>	
Slice commutes with the operator	<code>sliceAt_bellman, sliceAt_</code>	OLG/Continuity
	<code>augStageValue</code>	
Augmented stage policy	<code>augStagePolicy, augStagePolicy_eq</code>	OLG/Continuity
Jointly continuous stage policy	<code>continuousOn_augStagePolicy</code>	OLG/Continuity
Cohort recursion in the augmented state	<code>augCohortAssets, augCohortAssets_</code>	OLG/Continuity
	<code>eq</code>	
Theorem 6.3	<code>continuousOn_olgSupply</code>	OLG/Continuity
Theorem 6.4	<code>exists_olgEquilibrium</code>	OLG/Continuity
Saving rises iff consumption responds little	<code>stagePolicy_le_iff_consumption_le</code>	OLG/ RateMonotone
Terminal stage, unconditional	<code>stagePolicy_mono_zero</code>	OLG/ RateMonotone
Condition (1)	<code>StageSlack</code>	OLG/ RateMonotone
Theorem 7.1, step	<code>stagePolicy_mono_step</code>	OLG/ RateMonotone
Theorem 7.1, induction	<code>stagePolicy_mono_of_stageSlack</code>	OLG/ RateMonotone
Theorem 7.1, aggregate	<code>olgCapital_le_of_stageSlack</code>	OLG/ RateMonotone
Human wealth, price of the path	<code>humanWealth, pathPrice</code>	OLG/Certainty Equivalent
Closed form $c_0 = W/D$	<code>consumption_mul_pathPrice</code>	OLG/Certainty Equivalent
Saving rises iff the cross product falls	<code>consumption_antitone_iff</code>	OLG/Certainty Equivalent

Result	Lean name	Module
Theorem 7.2	<code>consumption_antitone_of_le_one</code>	OLG/Certainty Equivalent
Same household, different limit	<code>withFloor</code>	OLG/Borrowing Limit
Stages at a general limit	<code>finiteValue,</code> <code>finitePolicy,</code> <code>finiteConsumption</code>	OLG/Borrowing Limit
Theorem 7.4, consumption	<code>finiteConsumption_le_of_floor_le</code>	OLG/Borrowing Limit
Theorem 7.4, saving	<code>le_finitePolicy_of_floor_le</code>	OLG/Borrowing Limit
Theorem 7.5	<code>finiteConsumption_le_shift</code>	OLG/Borrowing Limit
Terminal slope is R	<code>stageConsumption_sub_zero</code>	OLG/ MPCBound
Theorem 7.6	<code>stageConsumption_sub_le_of_slope</code>	OLG/ MPCBound
Guaranteed human wealth	<code>minHumanWealth, lowMPC</code>	OLG/Lower Sandwich
The corner is respected	<code>lowMPC_mul_minHumanWealth_le</code>	OLG/Lower Sandwich
Theorem 7.7	<code>minHumanWealth_</code> <code>le_</code> <code>stageConsumption</code>	OLG/Lower Sandwich
Power mean under a shift	<code>powerMean_add_const_le</code>	OLG/Incidence
Risk-adjusted, capped human wealth	<code>riskHumanWealth</code>	OLG/Incidence
Theorem 7.8	<code>riskHumanWealth_</code> <code>le_</code> <code>stageConsumption</code>	OLG/Incidence
Theorem 7.9	<code>cap_necessary_of_corner</code>	OLG/Incidence
Definition 6.5	<code>StageRegions, allRegions</code>	OLG/LifeCycle
Sandwich on a reachable family	<code>stageConsumption_</code> <code>le_</code> <code>stageHumanWealth_on</code>	OLG/LifeCycle
Floor on a reachable family	<code>le_stagePolicy_on</code>	OLG/ Equilibrium
Crude reachable bound	<code>reach, reach_mono</code>	OLG/Reachable
Proposition 6.6	<code>reachRegions</code>	OLG/Reachable
Sandwich along a cohort's path	<code>stageConsumption_</code> <code>le_</code> <code>stageHumanWealth_reach</code>	OLG/Reachable
Floor along a cohort's path	<code>le_stagePolicy_reach</code>	OLG/ Equilibrium
Floor at the household's own state	<code>le_stagePolicy_state_on</code>	OLG/ Equilibrium
Proposition 6.13	<code>olgCapital_le_of_invariant</code>	OLG/Existence
Affine coefficients	<code>cohortSlope, cohortFloor</code>	OLG/Existence
Proposition 6.10	<code>cohortFloor_</code> <code>add_</code> <code>mul_</code> <code>le_</code> <code>cohortAssets, le_olgCapital</code>	OLG/Existence
Mean human wealth recursion	<code>sum_mul_stageHumanWealth_succ</code>	OLG/Existence
Theorem 6.11	<code>stageMPC_</code> <code>mul_</code> <code>sum_</code> <code>stageHumanWealth_le</code>	OLG/Existence

Result	Lean name	Module
The aggregate floor climbs	sum_mul_cohortFloor_mono, sum_mul_cohortFloor_nonneg	OLG/Existence
Theorem 6.16	exists_olgEquilibrium_of_bounds, exists_olgEquilibrium_rate	OLG/Existence
Proposition 6.7	stageMPC_mul_mpcSum, stageMPC_mul_one_sub_pow	OLG/LifeCycle
The geometric sum is bounded	sub_mul_mpcSum_le_one	OLG/LifeCycle
Constant-coefficient slope	cohortSlope_mul_mpcSum_succ	OLG/Existence
Proposition 6.8	mul_geomSum_le_cohortSlope	OLG/Existence
Human wealth in closed form	sum_mul_stageHumanWealth_eq	OLG/Existence
The floor in closed form	geomFloor_le_sum_mul_cohortFloor	OLG/Existence
Theorem 6.19	exists_olgEquilibrium_of_geomBounds	OLG/Existence
Definition 8.1	ValueGap	OLG/Rate Monotone
Theorem 8.2	stagePolicy_mono_of_valueGap, stagePolicy_mono_of_valueGap_all	OLG/Rate Monotone
Capital from the value gap	olgCapital_le_of_valueGap	OLG/Rate Monotone
The guaranteed asset path	assetFloor, assetFloor_le_stagePolicy	OLG/Asset Incidence
Human wealth with assets in it	assetHumanWealth, assetHumanWealth_mono	OLG/Asset Incidence
Theorem 7.10	assetHumanWealth_le_stageConsumption	OLG/Asset Incidence
It dominates the asset-free bound	riskHumanWealth_le_assetHumanWealth	OLG/Asset Incidence
Proposition 6.14	cohortAssets_le_of_incidence, olgCapital_le_of_incidence	OLG/Existence
Existence with the sharper ceiling	exists_olgEquilibrium_of_sharpBounds	OLG/Existence
Proposition 6.15	cohortCeil, cohortAssets_le_cohortCeil	OLG/Existence
Ceiling on aggregate capital	olgCapital_le_cohortCeil	OLG/Existence
Theorem 6.17	exists_olgEquilibrium_affine	OLG/Existence
Proposition 6.12	cohortFloorAt, cohortFloorAt_le_cohortAssets	OLG/Existence
Floor on capital, clamped	le_olgCapital_at	OLG/Existence
Theorem 6.18	exists_olgEquilibrium_clamped	OLG/Existence
The summed pairing	summedPairing, SummedLasryLionsMonotone	Equilibrium/LasryLions
Theorem 9.1	summedPairing_nonpos_of_pathEquilibrium	Equilibrium/LasryLions
Date-by-date is a special case	summedLasryLionsMonotone_of_forall, summedPairing_eq_zero_of_pathEquilibrium	Equilibrium/LasryLions
Theorem 9.2	aggregate_eq_of_pathEquilibrium	Equilibrium/LasryLions

Result	Lean name	Module
Costs seeing the whole family	famPathCost, IsFamPathEquilibrium,	Equilibrium/
	famSummedPairing	LasryLions
The engine in that generality	famSummedPairing_ nonpos_ of_	Equilibrium/
	pathEquilibrium	LasryLions
Theorem 9.3	famLifetime, famSummedPairing_	Equilibrium/
	scalar, statistic_ eq_ of_	LasryLions
	increasingDifferences	
Proposition 9.5	famCost, famSummedPairing_eq_neg_	Equilibrium/
	losses	LasryLions
Each cross-loss is nonnegative	famCost_ le_ of_ pathEquilibrium,	Equilibrium/
	famSummedPairing_nonpos_of_losses	LasryLions
Proposition 9.6	monotone_ of_ comonotone_ of_	Equilibrium/
	strictMono	LasryLions
The accumulated gap	gapAccum, gapAccum_nonneg	OLG/Increment
Theorem 10.1	negPart_le_gapAccum, neg_gapAccum_	OLG/Increment
	le	
Aggregate form, single crossing	sum_neg_gapAccum_le, lt_ of_	OLG/Increment
	increment_bound	
Superadditive certainty equivalent	powerMean_add_const_le	OLG/Incidence
Theorem 10.2	stageMPC_mul_le_stageConsumption_	OLG/Propensity
	sub	Floor
Theorem 10.3	cohortAssets_sub_le_cohortSlope	OLG/Propensity
		Floor
Theorem 10.4	incomeStep, stageStep_of_incomeFn,	OLG/Mass
	sum_mul_incomeStep	Weighted
The marginal at every age	sum_mul_incomeStep_iterate, sum_	OLG/Mass
	mul_le_of_le	Weighted